

Action-Angle Clustering of Stellar Streams in Merging Galaxies

MUQING WANG¹ AND TJITSKE STARKENBURG¹

¹*Northwestern University*

ABSTRACT

Stellar streams are sensitive tracers of the Galactic gravitational potential and therefore provide a promising probe of dark matter structure in the Milky Way halo. Many stream-based potential constraints rely on the expectation that debris from a cold progenitor remains clustered in action space when evaluated in the correct potential. However, the Milky Way is neither isolated nor static, and merger-induced perturbations may alter both the stream population and the reliability of action-space clustering as a potential diagnostic.

In this work, we test the robustness of action-space clustering using simulated globular-cluster streams evolved in live Milky Way-like halos undergoing mergers of different mass ratios and orbital configurations. We identify locally coherent stream cohorts using observable phase-space criteria, including smooth along-stream velocity profiles, unfolded phase-space structure, limited spatial extent, and locations away from orbital turning points. We then fit spherical NFW potentials by maximizing the one-dimensional Kullback–Leibler divergence (KLD) of the radial-action distribution and compare the fitted action clouds with those obtained in the frozen-snapshot Basis Function Expansion reference potential.

At the distance-matched fiducial 1:10 merger snapshot, the selected cohorts recover a useful inner-halo enclosed-mass profile. Across the full time evolution, however, stronger mergers remove usable coherent material earlier and produce larger and more rapidly varying mass biases, whereas the 1:100 merger remains comparatively stable. We identify three distinct limitations: loss of coherent tracer material, failure of the KLD-maximized potential to reproduce the correct instantaneous action cloud, and a local-to-global mismatch caused by halo non-sphericity. While usable cohorts remain, the largest mass errors are generally associated with the action-cloud mismatch, whereas the local-to-global contribution remains mostly below ten percent even in the 1:2 merger. These results show that observable cohort selection can improve the robustness of action-clustering fits, but cannot by itself guarantee unbiased potential recovery in a strongly perturbed halo.

Keywords: Galaxies (573) —

1. INTRODUCTION

Stellar streams form through the tidal disruption of globular clusters and dwarf galaxies. Because their stars originate from dynamically cold progenitors and subsequently follow neighboring trajectories in the Galactic gravitational field, stream debris often occupies a relatively compact region of phase space. The phase-space structure of the stellar stream carry valuable information about the mass profile, shape, and time evolution of the Milky Way halo.

Constraining the mass profile of the Milky Way is a central goal of near-field cosmology, since it constrains both the distribution of dark matter and the dynamical history of our galaxy. One promising way to extract this information is through the clustering of the stream stars in the action space. In a time-independent and integrable potential, the actions of individual stars are conserved. Debris stripped from a cold progenitor is

therefore expected to remain localized in action space when evaluated in the correct gravitational potential. This motivates potential-fitting methods that search for the model in which stream stars appear maximally clustered in action space. [R. E. Sanderson et al. \(2015\)](#) formalized this idea using the Kullback–Leibler divergence (KLD) as a quantitative measure of action-space clustering, and subsequent work like (e.g. [S. Reino et al. 2021, 2022](#)) applied and tested this approach using both mock and observed stream samples. These studies demonstrate that action-space clustering can provide strong constraints on the enclosed mass of the Milky Way, especially when multiple streams or carefully chosen stream regions are used.

However, the success of action-space clustering relies on assumptions that are only approximately satisfied in the real Milky Way. The Galactic halo is not an isolated, static, spherical system. The Large Magellanic Cloud

(LMC), the Small Magellanic Cloud, and the Sagittarius dwarf galaxy perturb both the shape and the kinematics of the Milky Way halo. In particular, the LMC can induce a large-scale response of the Galactic halo and can directly perturb stellar streams, changing their orbital tracks, velocities, and their constraining power (e.g. C. Vera-Ciro & A. Helmi 2013; D. Erkal et al. 2018; N. Shipp et al. 2019; R. A. N. Brooks et al. 2024, 2025). These effects raise an important question: when the true potential is time-dependent and non-spherical, does maximizing action-space clustering still recover the correct underlying mass profile?

Recent work has begun to address this question. A. Arora et al. (2022) showed that the stability of action-space clustering depends on the time evolution of the host potential and on the merger history of the galaxy. R. A. N. Brooks et al. (2024, 2025) further demonstrated that the deforming Milky Way halo induced by the LMC can affect the recovery of the spherically averaged Milky Way mass profile from stream actions. These results show that action-space clustering remains a promising tool, but its reliability must be tested in realistic, perturbed halos. They also suggest that not all stream material is equally useful: the constraining power of a stream depends on its phase-space coherence, orbital phase, and degree of perturbation.

In this work, we test the robustness of action-space clustering using simulated globular-cluster streams evolved in Milky Way-like halos undergoing mergers of different mass ratios and orbital configurations. Rather than using all available stream material, we select local coherent cohorts with smooth along-stream velocity profiles, unfolded phase-space structure, limited spatial extent, and locations away from radial turning points. These criteria are designed to isolate stream regions whose present-day kinematics provide evidence that their instantaneous action distributions may remain comparatively compact. For each selected cohort, we fit a spherical NFW potential by maximizing the one-dimensional KLD of its radial-action distribution. We compare the resulting fitted action clouds with those obtained in the frozen-snapshot BFE reference potential and quantify the corresponding enclosed-mass bias. We further decompose this bias into a fit-to-local component, which measures whether the optimized spherical potential reproduces the gravitational field sampled by the cohorts, and a local-to-global component, which measures whether that locally sampled field represents the globally spherical-averaged potential.

Our results show that stronger mergers generally remove usable coherent material earlier and produce larger and more rapidly varying mass biases, while weaker mergers preserve useful potential constraints for longer. We identify three distinct limitations: the disappearance of coherent tracer material, failure of the KLD-maximized model to reproduce the correct instantaneous action-cloud structure, and local-to-global mis-

match in a non-spherical halo. The second failure mode is generally responsible for the largest mass errors while usable cohorts remain, whereas the local-to-global correction is typically smaller.

Section 2 introduces the theoretical action-space and KLD framework. Section 3 presents the observable cohort-selection criteria. Section 4 describes the simulations, examines the time-dependent mass recovery, and three failure modes. Section 5 discusses the dependence on perturber orbital configuration, the limitations of observable cohort selection, and possible directions for future work.

2. THEORETICAL FRAMEWORK

2.1. Canonical Actions and Potential-dependent Action Clustering

For a Hamiltonian system with n degrees of freedom, the canonical action variables are well defined for regular motion confined to invariant tori. In particular, Liouville integrability requires n functionally independent integrals of motion F_i that are mutually in involution,

$$\{F_i, F_j\} = 0 \quad \text{for all } i, j. \quad (1)$$

When the corresponding common level set is regular, compact, and connected, it forms an invariant n -torus. The canonical actions are then defined by integrals over the n independent fundamental cycles γ_i of this torus:

$$J_i = \frac{1}{2\pi} \oint_{\gamma_i} \mathbf{p} \cdot d\mathbf{q}. \quad (2)$$

In separable canonical coordinates, this expression reduces to

$$J_i = \frac{1}{2\pi} \oint_{\gamma_i} p_i dq_i. \quad (3)$$

Here, q_i and p_i are canonically conjugate variables. The actions label the invariant torus and remain constant under the integrable Hamiltonian, while the conjugate angles evolve along it.

A static spherically symmetric potential is integrable. Away from degenerate orbits, the energy H , total angular momentum squared L^2 , and azimuthal angular momentum L_z provide three functionally independent integrals in involution. Consequently, three canonical actions (J_r, J_θ, J_ϕ) can be defined. Throughout the following expressions, momenta and energies are taken per unit stellar mass. The energy is

$$\begin{aligned} E &= \frac{1}{2} \left(p_r^2 + \frac{p_\theta^2}{r^2} + \frac{p_\phi^2}{r^2 \sin^2 \theta} \right) + \Phi(r) \\ &= \frac{1}{2} \left(p_r^2 + \frac{L^2}{r^2} \right) + \Phi(r), \end{aligned} \quad (4)$$

where

$$L^2 = p_\theta^2 + \frac{p_\phi^2}{\sin^2 \theta}. \quad (5)$$

Since the azimuthal angle ϕ is cyclic in a spherical potential, its conjugate momentum is conserved:

$$p_\phi = L_z. \quad (6)$$

So that the azimuthal action is:

$$J_\phi = \frac{1}{2\pi} \oint p_\phi \cdot d\phi = p_\phi = L_z \quad (7)$$

Using the expression of angular momentum shown in Equ.4, the polar action can be expressed as:

$$J_\theta = \frac{1}{2\pi} \oint p_\theta \cdot d\theta = \frac{1}{\pi} \int_{\theta_{min}}^{\theta_{max}} \sqrt{L^2 - \frac{L_z^2}{\sin^2\theta}} \cdot d\theta \quad (8)$$

Using $\sin(\theta_{min}) = \sin(\theta_{max}) = \frac{|L_z|}{L}$, the integral gives:

$$J_\theta = L - |L_z| \quad (9)$$

Finally, using Equ.4, the radial action can also be acquired:

$$J_r = \frac{1}{\pi} \int_{r_{min}}^{r_{max}} \sqrt{2[E - V(r)] - \frac{L^2}{r^2}} \cdot dr \quad (10)$$

where r_{min} and r_{max} are the radial turning points in the adopted spherical potential. For the fitted static spherical halo models, we calculate these canonical actions using **Agama**. The code determines the allowed radial interval and evaluates the corresponding canonical actions.

Stars stripped from a dynamically cold progenitor begin on nearby orbits and therefore occupy a relatively compact region in action space when evaluated in an appropriate integrable potential. This compactness contains information about the potential. Stars from the same stream generally sample different orbital phases. *An appropriate potential maps these different phase-space points onto nearby invariant tori*, whereas an incorrect trial potential generally assigns phase-dependent action errors, broadening or distorting the inferred action cloud. This is the physical basis of action-clustering potential inference (R. E. Sanderson et al. 2015; S. Reino et al. 2021, 2022).

This statement is not an identity: a wrong potential may translate an action cloud without greatly broadening it, and a short stream segment may provide insufficient phase coverage to distinguish trial models. Potential recovery therefore requires both a dynamically cold tracer and sufficient diversity in orbital phase, radius, energy, or angular momentum. Combining multiple streams on different orbits can reduce individual-stream degeneracies and biases (R. E. Sanderson et al. 2015; S. Reino et al. 2021).

2.2. 1D Kullback-Leibler Divergence

We quantify radial-action clustering using the Kullback–Leibler divergence (KLD). For a radial-action probability density $p(J_r)$ and a reference density $q(J_r)$,

$$D_{KL}[p||q] = \int p(J_r) \ln \left[\frac{p(J_r)}{q(J_r)} \right] dJ_r. \quad (11)$$

where $p(J_r)$ is the probability density of the radial action distribution at the points J_r , and $q(J_r)$ is the background distribution.

Also, following the work by R. E. Sanderson et al. (2015), I use the Modified Breiman Density Estimator to calculate the density profile of the action points. According to B. J. Ferdosi et al. (2011), a fine grid will be built across the action space, and the density value at the center of each fine grid $J_{r,k}$ will be evaluated using the following formula:

$$p(J_{r,k}) = \frac{1}{N} \sum_{i=1}^N (\sigma \lambda_i)^{-d} K_E \left(\frac{J_{r,k} - J_{r,i}}{\sigma_r \lambda_i} \right) \quad (12)$$

Here, i is the star index, N stands for the total number of stars, and d stands for the dimensionality of the action space. This estimator uses the Epanechnikov Kernel K_E instead of a Gaussian Kernel:

$$K_E(t) = \begin{cases} \frac{d+2}{2V_d} (1-t^2), & \text{if } t^2 < 1, \\ 0, & \text{otherwise.} \end{cases} \quad (13)$$

, where V_d is the volume of the unit sphere in d -dimensional space. We can easily see the the kernel contribution decreases with normalized distance from the sample point to the grid center, which is as expected. Following the analysis in B. J. Ferdosi et al. (2011), the window sizes σ_j in each direction in the action space is derived from coordinate percentiles in each dimension in the action space:

$$\sigma_r = \min \left\{ \frac{P_{80}(J_{r,i}) - P_{20}(J_{r,i})}{\ln N} \right\} \quad (14)$$

, where $P_n(J_{r,i})$ is the n th percentile of the radial action values. In addition, define a pilot density $p_{\text{pilot}}(J_{r,i})$ to describe the degree of clustering for the data point by using Equ. 12, with $\lambda_i = 1$. And the adaptive bandwidth λ_i of this estimator is designed to be inversely proportional to the $p_{\text{pilot}}(J_{r,i})$:

$$\lambda_i = \left(\frac{p_{\text{pilot}}(J_{r,i})}{g} \right)^{-\alpha} \quad (15)$$

, where g is the geometric mean of the pilot density:

$$g = \exp \left[\frac{1}{N} \sum_{i=1}^N \ln p_{\text{pilot}}(J_{r,i}) \right] \quad (16)$$

and $\alpha = \frac{1}{d}$ is the adaptive-bandwidth exponent. This adaptive bandwidth shrinks kernels in dense regions and widens them in sparse regions, thus increasing the density value where the action points cluster tightly.

Also, for those grids with overly small $p(J_r)$ values, a floor value (like 10^{-9}) needs to be assigned to prevent error. Luckily, the function "modified_breiman_density" in the "MBEtree" package can realize the above calculation for $p(J_r)$.

Following the action-clustering framework of R. E. Sanderson et al. (2015) and its applications by S. Reino et al. (2021, 2022), trial gravitational potentials are compared through the degree of clustering of the actions calculated from the same observed phase-space data. In this work, the reference distribution $q(J_r)$ is uniform over one fixed and sufficiently wide action interval $[J_{r,\min}, J_{r,\max}]$ that is shared by every trial potential,

$$q(J_r) = \frac{1}{J_{r,\max} - J_{r,\min}}. \quad (17)$$

The interval $[J_{r,\min}, J_{r,\max}]$ is kept the same for all trial potentials. In practice, we choose this universal range to be sufficiently wide to contain the radial actions produced over the full set of potentials considered, including deliberately extreme choices of the potential parameters. This ensures that changes in the KLD are not driven by changing the support of the background distribution from one potential to another.

Thus, the KLD measures how strongly the radial-action distribution departs from a uniform distribution over the fixed reference range. Since the reference interval is held fixed for all trial potentials, maximizing D_{KL}^I is equivalent to finding the potential in which the stream is most concentrated in J_r , or equivalently, has the lowest radial-action entropy. A mismatched potential generally broadens the inferred J_r distribution of a stream, thereby reducing the KLD. The KLD test used here is based on the expectation that the most accurate potential produces the sharpest radial-action clustering relative to the same uniform background. The success of this method has been demonstrated in static test potentials and in applications combining multiple stellar streams (R. E. Sanderson et al. 2015; S. Reino et al. 2021, 2022).

2.3. Reference-action Evolution in a Merging System

A live merger creates a conceptually different problem. The actual Hamiltonian is time dependent and generally non-integrable, so the stars may not have globally defined canonical actions of the actual system. However, canonical actions of an integrable spherical reference model can still be assigned to each measured phase-space point. These reference actions are model-dependent coordinates, not invariants of the merger Hamiltonian.

Let the actual Hamiltonian be $H(\mathbf{z}, t)$ and consider a cohort-common, smoothly evolving spherical reference

Hamiltonian

$$H_0(\mathbf{z}; \boldsymbol{\lambda}(t)), \quad (18)$$

, in which $\mathbf{z}$ is the actual position of the star, and $\boldsymbol{\lambda}$ is the spherical potential parameter. The complete actual-minus-reference residual is

$$\delta H(\mathbf{z}, t; \boldsymbol{\lambda}) \equiv H(\mathbf{z}, t) - H_0(\mathbf{z}; \boldsymbol{\lambda}). \quad (19)$$

For star a , we evaluate the reference action at its actual phase-space position,

$$J_{0,i}^{(a)}(t) \equiv J_i[\mathbf{z}_a(t); \boldsymbol{\lambda}(t)]. \quad (20)$$

The actual orbit need not lie on an invariant torus of the merger potential. Equation 20 only assigns its current phase-space point to a torus of the instantaneous spherical reference model.

Taking the total derivative along the actual trajectory gives

$$\begin{aligned} \frac{dJ_{0,i}^{(a)}}{dt} &= \left\{ J_{0,i}^{(a)}, H \right\} + \dot{\lambda}^\alpha \frac{\partial J_{0,i}^{(a)}}{\partial \lambda^\alpha} \bigg|_{\mathbf{z}_a} \\ &= -\frac{\partial \delta H^{(a)}}{\partial \theta_{0,i}} + \dot{\lambda}^\alpha \frac{\partial J_{0,i}^{(a)}}{\partial \lambda^\alpha} \bigg|_{\mathbf{z}_a}. \end{aligned} \quad (21)$$

, where repeated reference-model parameter indices α are summed. The first term is the physical response of the reference action to the actual-minus-reference Hamiltonian. The second term is the remapping caused by the evolution of the reference potential itself. This identity is exact when the full residual Hamiltonian is retained, the reference-action map is differentiable, and all quantities are evaluated at the actual phase-space point in one common canonical frame. If the Hamiltonian residual is purely a potential difference, $\delta H = \delta \Phi(\mathbf{x}, t)$, then:

$$\frac{dJ_{0,i}^{(a)}}{dt} = \delta \mathbf{F}^{(a)} \cdot \frac{\partial \mathbf{x}^{(a)}}{\partial \theta_{0,i}} + \dot{\lambda}^\alpha \frac{\partial J_{0,i}^{(a)}}{\partial \lambda^\alpha} \bigg|_{\mathbf{z}_a}. \quad (22)$$

Here $\delta \mathbf{F}$ is the force residual caused by the potential difference, and $\partial \mathbf{x} / \partial \theta_{0,i}$ is a reference-torus tangent evaluated at the actual phase-space point.

For two cohort members a and b ,

$$\begin{aligned} \frac{d}{dt} (J_{0,i}^{(a)} - J_{0,i}^{(b)}) &= -\Delta_{ab} \left(\frac{\partial \delta H}{\partial \theta_{0,i}} \right) \\ &\quad + \dot{\lambda}^\alpha \Delta_{ab} \left(\frac{\partial J_{0,i}}{\partial \lambda^\alpha} \bigg|_{\mathbf{z}} \right), \end{aligned} \quad (23)$$

where $\Delta_{ab} X \equiv X^{(a)} - X^{(b)}$. If the member-to-member differences on the right-hand side remain small over an extended interval, the pairwise reference-action separations vary slowly, allowing the reference-action cloud to undergo approximately semi-rigid translation. If it

varies strongly among members, their pairwise action separations grow and the cloud diffuses or distorts.

It is useful to distinguish this merger-driven loss of coherence from the potential-identification mechanism of Section 2.1. The observed width of a trial reference-action cloud contains at least two conceptually different contributions:

1. *evolutionary decoherence*: differential physical response and differential reference-coordinate remapping accumulated during the merger broaden the cohort even under a physically plausible spherical reference history; and
2. *trial-potential mismatch*: at the current snapshot, an incorrect spherical trial model maps stars at different orbital phases inconsistently and introduces additional differential action shifts.

The first contribution contaminates the clustering signal, whereas the second is the source of the potential information used by the KLD fit. If evolutionary decoherence dominates, maximizing the KLD may no longer identify a useful spherical mass model. Time-dependent and non-adiabatic halo deformations are known to reduce or displace stream clustering in action space and can bias spherical mass recovery (A. Arora et al. 2022; R. A. N. Brooks et al. 2024).

In Section 3, a cohort-selection procedure will be introduced. It attempts to retain cohorts using which the first source of reference action cloud broadening will be suppressed, so that the second source about the trial-potential mismatch can remain measurable.

2.4. Frozen-snapshot Radial-cycle Action-proxy

On the other hand, if the actual local potential structure differs strongly from being integrable, it will act as an additional factor that declines the accuracy of the KLD potential fitting.

Suppose we know the exact potential model, then at snapshot t_* , we can integrate a model orbit from the measured phase-space position, along which the stellar energy remains constant. If $\tau_{p,k}$ and $\tau_{p,k+1}$ are successive pericentric passages of this frozen model orbit, we define

$$\tilde{J}_{r,\text{BFE}}(t_*) \equiv \frac{1}{2\pi} \int_{\tau_{p,k}}^{\tau_{p,k+1}} p_r(\tau) \dot{r}(\tau) d\tau = \frac{1}{2\pi} \int_{\text{radial cycle}} p_r dr. \quad (24)$$

In an integrable spherical potential this reduces to the canonical radial action. In the frozen, non-spherical merger potential, however, one radial oscillation need not be a fundamental cycle of an invariant three-torus. The quantity may depend on phase, cycle number, and radial-cycle definition. We therefore refer to it as a frozen-potential radial-cycle action proxy rather than a true canonical action.

This proxy does not supply the theoretical basis of the KLD fit and is not an observable input to the method. It

is used only as a simulation diagnostic of disagreement between the radial-cycle structure generated by the fitted spherical model and that generated by the frozen potential field.

3. SELECTING COHORTS OF STREAM STARS

The theoretical framework above identifies two requirements for successful potential inference. First, KLD maximization can constrain a spherical trial potential only when the reference-action cloud of the selected stream material has avoided strong internal distortion. According to Equation 23, this requires the member-to-member differences in both the physical-response term and the reference-coordinate-remapping term to remain sufficiently small, so that the reference-action cloud may have undergone approximately semi-rigid translation and retained a small internal width. Second, the actual potential structure over the region sampled by the cohort must remain sufficiently representable by the adopted spherical integrable model family. Otherwise, the potential that maximizes reference-action clustering need not provide an accurate approximation to the locally sampled physical field.

In a real observation, neither the true time-dependent field nor the historical reference-action evolution is known. We therefore use only present-day phase-space information to select cohorts whose morphology and kinematics provide empirical evidence against strong accumulated differential mixing. These criteria improve the probability that a cohort is informative; they do not guarantee either historical semi-rigid translation or unbiased potential recovery.

Firstly, if the cohort members have experienced similar residual forces $\delta\mathbf{F}$ relative to a nearby spherical reference model, their accumulated velocity responses should also be similar. This favors a smooth velocity distribution along the stream. Since the gravitational potential and force used here are defined per unit stellar mass, the perturbation-induced velocity change may be written as

$$\delta\mathbf{v}(t_1 \rightarrow t_0) = \int_{t_1}^{t_0} \delta\mathbf{F}(t) dt + \text{const.}, \quad (25)$$

where t_0 is the present time and t_1 is an earlier time after the stream has settled into its orbit. A smooth present-day velocity profile does not prove that the residual force was coherent throughout this interval. Nevertheless, strong differential forcing accumulated over the stream history would generally tend to produce irregular or rapidly varying velocity differences. Present-day velocity smoothness can therefore be treated as a dynamical fossil that makes a historically coherent physical-response term in Equation ?? more plausible.

To implement this criterion, we introduce an along-stream coordinate s and project the position and velocity onto the local stream direction. We then fit a smooth

Illustration of the cohort-selection rules in radial phase space
ratio 1:10 | radial range 30-40 kpc




Figure 1. Illustration of the cohort-selection procedure in radial phase space for the 1:10 merger in the $30 \leq r \leq 40$ kpc radial range. **Left:** an arbitrary set of parent streams before cohort selection, including extended and folded phase-space structures. **Right:** coherent stream segments retained after requiring a smooth along-stream velocity profile, an unfolded monotonic phase-space structure, limited spatial extent, and location away from orbital turning points. The selection removes most of the extended and kinematically irregular stream material and retains short, locally coherent branches that are expected to provide tighter instantaneous-action distributions.

quadratic trend to the along-stream velocity:

$$v_{\text{fit}}(s) = as^2 + bs + c. \quad (26)$$

We retain only cohorts whose along-stream velocities deviate weakly from this fitted trend:

$$\Delta v = \frac{\sqrt{\langle [v(s) - v_{\text{fit}}(s)]^2 \rangle}}{\text{std}[v(s)]} < 0.1. \quad (27)$$

Secondly, we require an unfolded, single-valued phase-space distribution. Strong differential mixing tends to stretch and fold a stream, allowing stars at similar spatial positions to occupy different orbital branches with substantially different velocities and orbital phases. Removing such folded structures therefore makes it more likely that the selected members lie on one smooth branch of the reference-orbit mapping.

This criterion is also motivated by both factors appearing in Equation 23. Members on the same smooth phase-space branch are more likely to have similar reference-orbit tangents,

$$\frac{\partial \mathbf{x}}{\partial \theta_{0,i}},$$

and similar sensitivities of their reference actions to changes in the reference-potential parameters,

$$\frac{\partial J_{0,i}}{\partial \lambda^\alpha} \Big|_{\mathbf{z}}.$$

By contrast, wrapped branches may place stars at nearby spatial positions but with different velocities and radial phases, giving substantially different reference-orbit tangents and action sensitivities. The absence of wrapping therefore makes both the physical-response term and the reference-remapping term more likely to remain coherent across the cohort. Operationally, we retain only cohorts for which the fitted velocity trend $v_{\text{fit}}(s)$ is strictly monotonic over the selected interval.

Thirdly, we require the selected cohort to have a short spatial extent. Provided that the actual and reference force fields vary smoothly over the sampled region, a short segment is more likely to experience a small member-to-member variation in the residual force $\delta \mathbf{F}$. It also limits the spatial variation of the reference-orbit tangent and of the action sensitivity

$$\frac{\partial J_{0,i}}{\partial \lambda^\alpha} \Big|_{\mathbf{z}}.$$

This criterion therefore makes strongly differential physical response and strongly differential reference-coordinate remapping less probable at the current snapshot. When combined with velocity smoothness and unfoldedness, it also favors locally regular stream regions that may be described meaningfully by a common spherical reference model. It does not, however, demonstrate that the absolute perturbation is small or that the true local potential is exactly spherical.

At the same time, we impose a minimum radial span of 2 kpc:

$$\Delta r_{\text{cohort}} \equiv \max_a(r_a) - \min_a(r_a) \geq 2 \text{ kpc}, \quad (28)$$

where a indexes the stellar members of the cohort. This lower limit is an identifiability safeguard rather than a coherence requirement. A cohort spanning an extremely narrow radial interval may provide insufficient radial-phase and force-profile leverage to distinguish among spherical trial potentials, even when its inferred radial-action distribution is compact. The lower and upper extent limits therefore balance potential-identification leverage against local dynamical coherence. In the present-day sample analyzed here, all candidate cohorts already exceed the 2 kpc lower limit, so this requirement does not remove any additional present-day cohorts.

Fourthly, we exclude cohorts close to radial turning points. Near a pericentre or apocentre, v_r approaches zero and the local mapping between position and radial phase becomes poorly conditioned. Stars at similar radii may lie on the incoming and outgoing branches of the orbit and may consequently have different reference-orbit tangents and different sensitivities to changes in the reference-potential parameters. We therefore require the selected members to remain away from radial turning points by imposing a threshold on their radial velocities. This requirement is not redundant with the unfoldedness criterion: a phase-space branch may remain unfolded while still lying close to a radial turning point.

To conclude, I use the following 4 criteria to act as a pass for selecting the coherent stellar cohorts:

- i) Smooth Velocity Profile, required by Equ. 27;
- ii) Unfolded Phase-space Distribution, enforced by the monotonicity of $v_{fit}(s)$;
- iii) Short Spatial Extent;
- iv) Away from Orbital Turning Points, enforced by the threshold in $|v_r|$.

These four present-day criteria provide observable evidence against strong accumulated differential mixing, strongly differential reference-action evolution, and severe local dynamical irregularity. Their present-day satisfaction is interpreted as a dynamical fossil that makes a previous approximately semi-rigid translation of the reference-action cloud more probable. The criteria do not prove such a history, nor do they guarantee that the true local potential is accurately represented by a spherical model. Rather, they are designed to select stream material for which the two principal causes of KLD-fitting failure—evolutionary distortion of the reference-action cloud and a severe mismatch between the locally sampled dynamics and the spherical trial family—are less likely to dominate.

Fig. 1 illustrates how strongly these criteria reduce the parent-stream population to a small set of locally coherent segments. Before selection, the radial phase-space distribution contains extended streams with substantial

curvature, overlap, and folded structures, including material close to orbital turning points. After applying the four selection criteria, the retained cohorts occupy much shorter and more nearly single-valued branches in radial phase space. The selected segments therefore exclude much of the spatially extended, kinematically irregular, and strongly phase-mixed material that would otherwise broaden the inferred action distribution.

This comparison also emphasizes that the selection is deliberately restrictive rather than simply removing a small number of outlying stars. The resulting cohorts represent only those local portions of the streams whose present-day phase-space structure provides strong evidence for coherent evolution history.

4. COMPUTATIONAL WORKS

4.1. Simulation of merging halos with stellar streams

We use the stellar-stream simulations from *Dancing Streams in Merging Halos* (Weerasooriya et al. 2025), a controlled suite designed to study how globular-cluster streams respond to a MW-LMC-like merger. The simulations first generate 1024 mock globular-cluster streams in a static MW-like Hernquist halo using the Fardal et al. (2015) mock-stream distribution function as implemented in Gala. Each stream contains 1000 stellar particles. The stream grid spans two different stellar masses, 10^4 and $10^6 M_\odot$, eight apocenter radii from 30 to 100 kpc, eight circularities from $\eta = 0.5$ to 1.0, and eight stream ages from 0.5 to 6 Gyr, giving 1024 streams in total.

After generation, the same stream population is inserted into live N-body halos and evolved with Gadget4. The host halo has mass $M_{\text{halo}} = 1.57 \times 10^{12} M_\odot$ and Hernquist scale radius $a = 40.85 \text{ kpc}$. The published *Dancing Streams in Merging Halos* suite describes the evolution of the streams in an isolated host and in a 1:5 MW-LMC-like merger, including fiducial, more circular, and radial perturber orbits. In this work, I also use additional related merger data products that extend the stream kinematics to a wider range of perturber-to-host mass ratios, 1:1, 1:2, 1:5, 1:10, and 1:100, with circular and radial perturber-orbit variants available for most ratios. This extended set allows me to study separately the effects of merger strength and perturber orbital geometry on the survival of coherent stream cohorts and on the bias of KLD-based spherical-NFW potential fits. The full simulations are evolved for 6 Gyr with 241 snapshots. In this work, we analyze snapshots 000–100, corresponding to 0–2.56 Gyr. Our main KLD-fitting analysis focuses on selected cohorts in fixed present-day radial ranges, especially 20–30 kpc, with additional comparisons in 30–40 kpc and 50–70 kpc.

The potential model of the merger halo data packs can be described by the Basis Function Expansion (BFE), which expresses the mass density distribution as a su-

perposition of basis functions:

$$\rho(\mathbf{x}, t) \approx \sum_{n, \ell, m} a_{n, \ell, m}(t) \rho_{n \ell m}(\mathbf{x}). \quad (29)$$

, where $\rho_{n, \ell, m}(\vec{x})$ is the basis function that consists of both the radial distribution and the angular distribution:

$$\rho_{n \ell m}(\mathbf{x}) = u_{n \ell}(r) Y_{\ell m}(\theta, \phi). \quad (30)$$

The relationship between the mass density distribution and the potential distribution can be described by Poisson's Equation:

$$\nabla^2 \phi(r) = 4\pi G \rho(r) \quad (31)$$

For those coefficients $a_{n, \ell, m}(t)$, I adopted the results worked out by David Amir Gonzalez, who is a graduate student at the University of Pennsylvania currently working with Professor Robyn Sanderson, upon their agreement. Their result accurately represent the later-stage potential structure in the simulation.

4.2. Present-day recovery and time-dependent inner-halo performance

4.2.1. Implementing the Selection Criteria

To test those selection criteria, I utilize them to pick coherent cohorts from the tidal streams simulated under an LMC-like merger, which has 1/10 of the mass of the Milky Way halo, merging into the Milky Way. The Milky Way halo is assumed to have a potential described by a spherical NFW model initially. Because the LMC is currently located at a Galactocentric distance of approximately 49.5 kpc (X. Wu et al. 2008) and is likely on its first infall or first passage about the Milky Way (G. Besla et al. 2007; N. Kallivayalil et al. 2013), we identify snapshot 55 ($t = 1.406$ Gyr), at which the simulated host-perturber separation is closest to the observed present-day MW-LMC separation, as the distance-matched present-day analogue used to examine the cohort-selection criteria.




Figure 2. Plot for host-perturber distance for the simulation of tidal streams in the Milky Way under an LMC-like merger. Snapshot 55 ($t = 1.406$ Gyr) is identified as the distance-matched present-day reference snapshot.

In general, the gravitational potential of a perturbed, triaxial MW halo should be described by a general expansion, like the Basis Function Expansion (BFE). However, since I use data from a simulation that assumes a spherical potential initially, and that fitting a general BFE requires a much longer running time along with more stellar members to constrain the parameters, I followed the work of Brooks et al and optimized the parameters of a spherical NFW potential model.

Using those criteria, I selected coherent cohorts, each with at least 100 stars, in every 10-kpc-wide radial range from the galactic center outward. Each coherent cohort will give an optimized NFW model, using which I can integrate for an enclosed mass profile from the galactic center outward. One thing worth noting is that the perturbation force varies so strongly in the region near the center that no coherent cohort can be picked in the region of $r < 10$ kpc.

The spherical enclosed mass implied by an arbitrary gravitational potential can be obtained from the radial gravitational acceleration on a spherical surface. The gravitational acceleration is

$$\mathbf{g}(\mathbf{x}, t) = -\nabla \Phi(\mathbf{x}, t), \quad (32)$$

while Poisson's equation gives

$$\nabla^2 \Phi(\mathbf{x}, t) = 4\pi G \rho(\mathbf{x}, t). \quad (33)$$

Combining Equations 32 and 33, we obtain

$$\nabla \cdot \mathbf{g}(\mathbf{x}, t) = -4\pi G \rho(\mathbf{x}, t). \quad (34)$$

We integrate Equation 34 over the volume V_r enclosed by a sphere of radius r :

$$\begin{aligned} \int_{V_r} \nabla \cdot \mathbf{g} d^3\mathbf{x} &= -4\pi G \int_{V_r} \rho(\mathbf{x}, t) d^3\mathbf{x} \\ &= -4\pi G M_{\text{sph}}(< r, t). \end{aligned} \quad (35)$$

Applying the divergence theorem gives

$$\oint_{S_r} \mathbf{g} \cdot d\mathbf{S} = -4\pi G M_{\text{sph}}(< r, t), \quad (36)$$

where S_r is the spherical surface bounding V_r .

On a sphere of radius r , the outward surface element is

$$d\mathbf{S} = r^2 \hat{\mathbf{r}} d\Omega, \quad (37)$$

and the radial component of the gravitational acceleration is

$$g_r(r, \Omega, t) \equiv \mathbf{g}(r, \Omega, t) \cdot \hat{\mathbf{r}}. \quad (38)$$

Also, we can define the angular average of the radial acceleration at fixed radius as

$$\langle g_r(r, \Omega, t) \rangle_\Omega \equiv \frac{1}{4\pi} \int_{4\pi} g_r(r, \Omega, t) d\Omega. \quad (39)$$

Substituting Equations 37, 38 and 39 into Equation 36 yields the spherically averaged enclosed mass:

$$M_{\text{sph}}(< r, t) = -\frac{r^2}{G} \langle g_r(r, \Omega, t) \rangle_\Omega. \quad (40)$$

The above equation does not require the gravitational potential to be spherical. It follows from the total gravitational flux through a spherical surface and therefore returns the mass enclosed by that surface implied by the potential model.

In my numerical calculation, the angular average in Equation 40 is evaluated numerically using $N_\Omega = 4096$ nearly uniformly distributed Fibonacci-sphere directions. Thus,

$$\langle g_r(r, \Omega, t) \rangle_\Omega \simeq \frac{1}{N_\Omega} \sum_{a=1}^{N_\Omega} g_r(r, \Omega_a, t), \quad (41)$$

and the enclosed mass is estimated as

$$M_{\text{sph}}(< r, t) \simeq -\frac{r^2}{G N_\Omega} \sum_{a=1}^{N_\Omega} g_r(r, \Omega_a, t). \quad (42)$$

The Fibonacci construction provides nearly equal-area sampling over the sphere and is sufficiently dense relative to the angular resolution of the NFW spherical model.

In order to improve the precision, I decided to look at the weight-average enclosed-mass profile. In the left panel of Fig.3, I plotted the enclosed mass profile from each cohort only in the radial range where the cohort is

located. Then I chopped the whole radial range into 0.1-kpc-wide bins. At each bin r_k , I weight-averaged over the enclosed mass values fitted from each cohort i that covers this bin:

$$M_{\text{fit}}(< r_k, t) = \frac{\sum_{i \in \mathcal{I}_k} N_i M_{\text{sph}, i}(< r_k, t)}{\sum_{i \in \mathcal{I}_k} N_i}, \quad (43)$$

where $\mathcal{I}_k$ denotes the selected cohorts whose fitted profiles are evaluated at r_k , and N_i is the number of selected stars in cohort i . The right panel of Fig. 3 shows the weight-averaged enclosed mass value at each radial bin from the galactic center outward.

From this result, we can see that, despite the error caused by the model simplification, the streams selected by those criteria retrieve the enclosed mass profile reasonably precisely, especially in the range of $r \in [10, 30]$ kpc. Outside this radial range, the discrepancy arises as we go further out, which is expected. At snapshot 55, the center of mass of the merger is about 49.5 kpc away from the host halo. At such a radial distance or outward, a stream segment is more likely to survive as a coherent cohort if it lives in an area farther away from the perturber, where the local Dark Matter particles are more dominant by those from the host halo. Using the fitted local spherical potential at those lower-density and less-perturbed region to retrieve the spherically-averaged enclosed mass would naturally underestimate the global enclosed mass. So at the current merger stage of the LMC merger, the KLD-maximization method based on those selection criteria discussed above would work well to constrain the potential model inside the radial range sampled by the merger particles.

4.2.2. Time-dependent inner-halo performance

The distance-matched snapshot demonstrates that the selected cohorts can recover a useful enclosed-mass profile at a particular stage of the merger. However, this single-snapshot result does not establish how long the method remains reliable as the halo and stream population evolve. We therefore repeat the cohort selection and spherical-NFW fitting at every snapshot from 0 to 100, corresponding to $0 \leq t \lesssim 2.56$ Gyr.

We first consider cohorts occupying the radial interval $20 \leq r/\text{kpc} \leq 30$. For this inner-halo analysis, the reference enclosed mass is measured directly from the host and perturber dark-matter particles, rather than from the BFE potential.

To characterize the surviving stream sample independently of the tested potential, we define the six-dimensional phase-space compactness of cohort i as

$$C_i = \left[\frac{1}{N_i} \sum_{j=1}^{N_i} \|z_{ij} - \bar{z}_i\|^2 \right]^{1/2}, \quad (44)$$




Figure 3. Enclosed-mass recovery from the selected coherent cohorts at snapshot 055, corresponding to the distance-matched present-day configuration of the fiducial MW-LMC-like merger. **Left:** the KLD-maximized spherical-NFW enclosed-mass profile inferred independently from each selected cohort, shown only over the radial interval sampled by that cohort; the black curve shows the true dark-matter enclosed-mass profile for reference. **Right:** the cohort-size-weighted combination of the individual fitted profiles compared with the same reference profile. The selected cohorts reproduce the enclosed mass most closely in the inner radial range, while increasingly large discrepancies appear farther outward where the merger-induced departure from a spherical equilibrium model becomes more important.

where j is the star index, N_i is the number of selected stars in cohort i , and

$$\mathbf{z} = \left(\frac{x}{10 \text{ kpc}}, \frac{y}{10 \text{ kpc}}, \frac{z}{10 \text{ kpc}}, \frac{v_x}{100 \text{ km s}^{-1}}, \frac{v_y}{100 \text{ km s}^{-1}}, \frac{v_z}{100 \text{ km s}^{-1}} \right). \quad (45)$$

The selected-cohort-weighted compactness is

$$\langle C \rangle_w = \frac{\sum_i N_i C_i}{\sum_i N_i}. \quad (46)$$

At each radial grid point r_k within the 10-kpc-wide radial bin, the weighted fitted enclosed mass is

$$M_{\text{fit}}(< r_k, t) = \frac{\sum_{i \in \mathcal{I}_k} N_i M_{\text{sph}, i}(< r_k, t)}{\sum_{i \in \mathcal{I}_k} N_i}, \quad (47)$$

where $\mathcal{I}_k$ denotes the selected cohorts whose fitted profiles are evaluated at r_k . We quantify the signed enclosed-mass bias by

$$b_M(t) = \left\langle \frac{M_{\text{fit}}(< r_k, t) - M_{\text{true}}(< r_k, t)}{M_{\text{true}}(< r_k, t)} \right\rangle_k \times 100\%. \quad (48)$$

Figure 4 summarizes the resulting evolution calculated under the fiducial 1:10 merger. The first column shows the number of selected cohorts and their mean stellar membership, the second column shows $\langle C \rangle_w$, and the

third column shows b_M . The distance-matched present-day reference time is marked by a light-red vertical line. The compactness generally increases with time, although it does not exhibit a single threshold that separates successful and unsuccessful fits. In section 4.3, we will discuss another set of quantities that are indeed related to the fitting precision. So far, the clearest merger-dependent hierarchy instead appears in the survival of the selected sample and in the reliability of the recovered mass.

The 1:2 merger loses most of its usable cohorts earliest and develops large, rapidly varying mass biases. The 1:5 and 1:10 cases retain selected material for longer, but subsequently exhibit substantial mass fluctuations. The 1:100 merger remains comparatively stable over most of the analyzed interval. Thus, stronger mergers generally shorten the interval over which a sufficiently large and dynamically coherent stream sample remains available and over which the fitted mass remains precise.

Occasional late-time snapshots may nevertheless return a small value of b_M . Such isolated agreement does not by itself demonstrate restored reliability. Because the cohorts are selected independently at each snapshot, the late-time sample can be dominated by an atypical subset of cohorts free of those moderately-perturbed yet passable members. We therefore interpret the time limit of the method from sustained behavior over multiple




Figure 4. Time-dependent performance of the selected-cohort KLD fitting for merger mass ratios 1:2, 1:5, 1:10, and 1:100 in the inner-halo radial range. The light red vertical line marks the distance-matched present-day reference time defined using the fiducial 1:10 merger. **Left column:** number of selected coherent cohorts and their average stellar membership. **Middle column:** selected-cohort-weighted six-dimensional phase-space compactness $\langle C_{6D} \rangle_w$. **Right column:** signed enclosed-mass bias of the KLD-maximized spherical-NFW fit relative to the true dark-matter enclosed mass. Stronger mergers generally remove usable coherent stream material earlier and produce larger mass biases, whereas the 1:100 merger system remains comparatively stable over most of the analyzed interval. The phase-space compactness evolves with time but does not by itself define a universal threshold for successful potential recovery.

snapshots, rather than from individual near-zero mass residuals.

These inner-halo results establish the phenomenology of failure, but they do not distinguish why an available cohort produces an incorrect enclosed-mass estimate. We address this question in Section 4.3 using streams in the outer radial range $80 \leq r/\text{kpc} \leq 100$, restricted to snapshot intervals over which the available BFE representation provides a sufficiently reliable description of the simulated force field.

For each selected cohort, we compare the canonical radial-action distribution calculated in the KLD-optimized spherical reference potential with the frozen-BFE radial-cycle action-proxy distribution defined in Equation ???. This comparison diagnoses disagreement between the radial-cycle mappings generated by the optimized spherical model and by the frozen BFE potential. We combine this diagnostic with a decomposition of the enclosed-mass bias into fitted-to-local and local-to-global contributions.

4.3. Three failure modes of the KLD potential fit

For every selected star, let $J_{r,\text{fit}}$ denote the canonical radial action evaluated in the KLD-optimized spherical NFW potential, and let $\tilde{J}_{r,\text{BFE}}$ denote the frozen-BFE radial-cycle action proxy defined in Equation 24. We characterize the displacement between the two cohort distributions using

$$\Delta \text{median}(J_r) = \text{median}(J_{r,\text{fit}}) - \text{median}(\tilde{J}_{r,\text{BFE}}), \quad (49)$$

and characterize their difference in clustering using

$$\Delta D_{\text{KL}}(J_r) = D_{\text{KL}}[J_{r,\text{fit}}] - D_{\text{KL}}[\tilde{J}_{r,\text{BFE}}]. \quad (50)$$

The two KLD values are evaluated using the same fixed reference distribution and action interval.

We retain the signs of both quantities. The sign of $\Delta \text{median}(J_r)$ records the direction of the displacement between the spherical-reference action cloud and the frozen-BFE radial-cycle proxy cloud. The sign of $\Delta D_{\text{KL}}(J_r)$ indicates whether the optimized spherical model produces a more or less clustered distribution than the frozen-BFE radial-cycle mapping.

These quantities should not be interpreted as differences between two sets of exact actions. Only $J_{r,\text{fit}}$ is a canonical radial action. The BFE quantity is a frozen-potential radial-cycle proxy. Their discrepancy is therefore a diagnostic of disagreement between the radial-cycle structures generated by the two potentials, rather than a direct metric of the potential difference itself.

Figure 5 provides a direct comparison between the evolution of the selected stream sample, the fitted-versus-BFE action-cloud discrepancies, and the resulting enclosed-mass bias in the outer-halo radial range of $80 \leq r/\text{kpc} \leq 100$. Two general pattern is visible

across the merger sequence: Firstly, the action discrepancy profiles have similar patterns to that of the enclosed mass; Secondly, the weaker mergers tend to preserve both a closer agreement between the fitted and BFE action clouds and a more stable recovery of the enclosed mass, whereas the stronger mergers develop larger action-space discrepancies together with larger and more rapidly varying mass biases. To diagnose the origin of these discrepancies, it is useful to separate the total signed mass bias into physically distinct contributions.

At radial grid point r_k , the fitted enclosed mass remains the cohort-size-weighted quantity

$$M_{\text{fit}}(< r_k, t) = \frac{\sum_i N_i M_{\text{sph},i}(< r_k, t)}{\sum_i N_i}. \quad (51)$$

For the BFE reference, we define the local spherical-equivalent mass at direction Ω as

$$M_{\text{local}, \Omega}(< r_k, \Omega, t) = -\frac{r_k^2}{G} g_r(r_k, \Omega, t), \quad (52)$$

and evenly average this quantity over the spatial directions sampled by the selected material,

$$M_{\text{local}}(< r_k, t) = \frac{1}{N_{\Omega,k}} \sum_{a=1}^{N_{\Omega,k}} M_{\text{local}, \Omega}(< r_k, \Omega_a, t), \quad (53)$$

where $N_{\Omega,k}$ is the number of spacial directions in the angular area sampled by the selected cohorts at r_k .

We therefore define

$$b_{M,\text{fit} \rightarrow \text{local}}(t) = 100 \left\langle \frac{M_{\text{fit}}(< r_k, t) - M_{\text{local}}(< r_k, t)}{M_{\text{true}}(< r_k, t)} \right\rangle_k, \quad (54)$$

and

$$b_{M,\text{local} \rightarrow \text{true}}(t) = 100 \left\langle \frac{M_{\text{local}}(< r_k, t) - M_{\text{true}}(< r_k, t)}{M_{\text{true}}(< r_k, t)} \right\rangle_k \quad (55)$$

in the radial range of $80 \leq r/\text{kpc} \leq 100$. The radial-grid average $\langle \dots \rangle_k$ is an ordinary equal-bin average. The first quantity measures whether the KLD-maximized spherical potential reproduces the effective local gravitational field sampled by the selected stream material, whereas the second quantity measures whether that local field is representative of the globally spherical-averaged BFE potential field. Thus, when a usable selected-cohort population is available, the total signed mass bias can be written as

$$b_M(t) = b_{M,\text{fit} \rightarrow \text{local}}(t) + b_{M,\text{local} \rightarrow \text{true}}(t). \quad (56)$$

In the following subsections, we will discuss about those two components more specifically.




Figure 5. Comparison between the KLD-fitted spherical-NFW model and the frozen-snapshot BFE reference potential in the $80 \leq r/\text{kpc} \leq 100$ radial range, restricted to snapshot intervals in which the BFE provides a reliable global reference. **Left column:** availability and stellar membership of the selected coherent cohorts. **Middle column:** signed fitted-minus-BFE discrepancies in the median radial action, $\Delta \text{median}(J_r)$, and in the radial-action clustering measure, $\Delta \text{KLD}(J_r)$. **Right column:** signed enclosed-mass bias of the fitted spherical model relative to the BFE reference. The stronger merger cases develop substantially larger discrepancies between the fitted and BFE action clouds together with larger mass errors, while the 1:100 merger case remains comparatively close to the BFE action structure and global enclosed mass.

4.3.1. Failure mode I: loss of coherent stream material

The first failure mode occurs when the merger stretches, folds, heats, or otherwise perturbs the streams strongly enough that few or no stellar segments satisfy the coherence criteria. In this regime, the limitation is not primarily that the KLD optimizer selects an incorrect potential; rather, the data required for the optimization are no longer available. Similar to the previous

inner-range analysis, the fitted results for a 1:2 merger show discontinuity at earlier time, due to the earlier removal of usable coherent stream material.

4.3.2. Failure mode II: mismatch between the spherical-reference and frozen-BFE radial-cycle mappings

A misspecified spherical potential may produce a compact radial-action distribution while placing that distri-




Figure 6. Fit-to-local enclosed-mass bias, $b_{M, \text{fit} \rightarrow \text{local}}$, versus signed fitted-minus-BFE action-cloud clustering $\Delta D_{\text{KL}} = D_{\text{KL}, \text{fit}} - D_{\text{KL}, \text{BFE}}$ for the 1:2, 1:5, 1:10, and 1:100 mergers in the BFE-valid 80-100 kpc snapshot intervals. Different symbols denote separate contiguous BFE-valid snapshot intervals. Black lines show Theil-Sen fits and the gray regions show the 95% moving-block-bootstrap ranges. The 1:2 case exhibits only shallow scalar trends.

bution at the wrong action-space location or producing a different degree of clustering from the frozen-snapshot BFE reference. To isolate this failure mode from the separate local-to-global effect discussed in Section 4.3.3, we compare the action-cloud discrepancies specifically with $b_{M, \text{fit} \rightarrow \text{local}}$, rather than with the total mass bias b_M .

Figure 6 and 7 shows these two signed action-cloud diagnostics against $b_{M, \text{fit} \rightarrow \text{local}}$ for the different merger strengths. We summarize each relation using a Theil-Sen robust linear estimator. Because adjacent snapshots within a BFE-valid interval are temporally correlated, the uncertainty in the fitted slope is estimated using a moving-block bootstrap within each contiguous interval. These regressions should be interpreted as descriptive diagnostics rather than universal linear scaling relations.

From those result, a larger ΔD_{KL} is associated with a more negative fit-to-local mass bias. One possible interpretation is that the spherical NFW fit partially com-

pensates for merger-induced perturbations by selecting a potential that artificially compresses the inferred J_r distribution. This over-clustering appears to be associated with a shallower locally inferred force field and therefore an underestimated enclosed mass. Whereas a positive displacement of the fitted median J_r tends to accompany a more positive enclosed mass bias. A possible interpretation of the positive relation is that a fitted NFW model with a larger enclosed mass produces a deeper effective potential and changes the reconstructed radial orbit so that the same observed phase-space coordinates enclose a larger radial phase-space area. The resulting increase in the fitted median J_r would then accompany a positive local enclosed-mass bias.

The resulting trends vary strongly with merger strength. The 1:5 and 1:10 mergers show appreciable relations between the signed action-cloud discrepancies and $b_{M, \text{fit} \rightarrow \text{local}}$. The 1:100 system occupies a much nar-




Figure 7. Fit-to-local enclosed-mass bias, $b_{M, \text{fit} \rightarrow \text{local}}$, versus signed fitted-minus-BFE action-cloud displacement $\Delta \text{median}(J_r) = \text{median}(J_{r, \text{fit}}) - \text{median}(J_{r, \text{BFE}})$ for the 1:2, 1:5, 1:10, and 1:100 mergers in the BFE-valid 80-100 kpc snapshot intervals. The 1:2 case exhibits only shallow scalar trends.

rower range in both action discrepancy and mass bias, as expected for the least strongly perturbed halo.

The 1:2 merger case is special. As shown in Figure 6 and 7, its fitted slopes are comparatively shallow: approximately -2.8 for the ΔD_{KL} relation and $+4.1$ for the $\Delta \text{median}(J_r)/10^4$ relation, with correspondingly weak rank correlations. The weak scalar relations in the 1:2 merger likely indicate a breakdown of a one-dimensional mapping between radial action-cloud mismatch and local enclosed-mass bias. In this strongly perturbed regime, the fitted and BFE action distributions may differ in shape, skewness, or multimodality in ways that are not captured by either ΔD_{KL} or $\Delta \text{median}(J_r)/10^4$. Moreover, J_r depends on the force profile over the full reconstructed orbit, which might extend far inside and outside 80–100 kpc, while $b_{M, \text{fit} \rightarrow \text{local}}$ compares radial forces over the selected radial shell and directions. In the strongly perturbed 1:2 merger case, changes in the global radial profile or higher multipoles of the potential can strongly affect J_r without producing

a proportional change in the local spherical-equivalent mass. Rapid changes in the surviving cohort population and the combination of disconnected snapshot intervals may introduce additional scatter. Therefore, the shallow slopes do not necessarily indicate the absence of an action-space failure, but rather that this failure is inadequately represented by either scalar diagnostic alone.

4.3.3. Failure mode III: local-to-global mismatch in a non-spherical potential

Even if a spherical fitted potential accurately reproduced the gravitational field locally sampled by a selected cohort, the resulting spherical-equivalent enclosed mass would not necessarily equal the globally spherical-averaged enclosed mass.

For the full BFE force field, the globally spherical-averaged enclosed mass is

$$M_{\text{true}}(< r, t) = -\frac{r^2}{G} \langle g_r(r, \Omega, t) \rangle_{\Omega}. \quad (57)$$




Figure 8. Local-to-true contribution to the enclosed-mass bias, $b_{M, \text{local} \rightarrow \text{true}}$, evaluated over the outer-range cohorts in the BFE-valid snapshot intervals. Each curve compares the spherical-equivalent mass inferred from the local BFE potential in the space sampled by the selected cohorts. The mismatch remains mostly below 10% even for the 1:2 merger and is generally only a few percent for the weaker mergers. Thus, merger-induced non-sphericity contributes to the total mass bias but is typically subdominant to the discrepancy between the KLD-fitted potential and the locally sampled BFE field.

At an individual direction Ω , we define

$$M_{\text{local}, \Omega}(< r, \Omega, t) = -\frac{r^2}{G} g_r(r, \Omega, t). \quad (58)$$

This quantity is the spherical enclosed mass that would reproduce the true BFE radial acceleration at that location.

We evaluate $M_{\text{local}, \Omega}$ along the angular directions sampled by the selected outer-halo material and form the unweighted sampled mean M_{local} . Then the corresponding local-to-true contribution to the mass bias is

$$b_{M, \text{local} \rightarrow \text{true}}(t) = 100 \left\langle \frac{M_{\text{local}}(< r_k, t) - M_{\text{true}}(< r_k, t)}{M_{\text{true}}(< r_k, t)} \right\rangle_k. \quad (59)$$

Figure 8 shows the time evolution of $b_{M, \text{local} \rightarrow \text{true}}$ during the same BFE-valid intervals used for the action-cloud comparison. The local-to-true mismatch is noticeably larger for the stronger mergers, but its absolute magnitude remains comparatively modest. Even in the 1:2 merger, $b_{M, \text{local} \rightarrow \text{true}}$ remains predominantly within approximately ten percent of zero over the sampled intervals. The 1:5 and 1:10 mergers typically show only several-percent offsets, while the 1:100 case remains still closer to zero.

Non-sphericity therefore contributes a genuine additional uncertainty to the interpretation of a locally fitted spherical potential, but the present decomposition indicates that it is *not the dominant source of the large mass failures* seen in the strongly perturbed cases. In particular, the tens-of-percent excursions in the fitted mass cannot generally be produced by a local-to-true term whose magnitude is mostly below ten percent. Most of the large discrepancy must instead arise in $b_{M, \text{fit} \rightarrow \text{local}}$, consistent with failure of the KLD-maximized spherical potential to recover the appropriate instantaneous local action structure.

The second local-to-true term in Equation 56 can nevertheless affect the detailed snapshot-to-snapshot total bias. Its sign can either reinforce or partially cancel the contribution of the first fit-to-local term, so it may introduce additional scatter when the total b_M is compared directly with a scalar action-cloud diagnostic. Its role should therefore be regarded as a *secondary non-sphericity correction*, rather than as the principal explanation for the breakdown of the KLD fit in the 1:2 merger.

4.3.4. Combined reliability conditions

The KLD method therefore has no single universal failure threshold. Reliable recovery requires the simultaneous satisfaction of three conditions:

1. a sufficiently large population of dynamically coherent stream cohorts must remain available;

2. the KLD-maximized fitted potential must reproduce the relevant location and clustering structure of the instantaneous action cloud under the true frozen-snapshot potential; and
3. the local gravitational field sampled by the retained cohorts must be sufficiently representative of the global potential distribution.

The new mass-bias decomposition shows that these three limitations are not equally important in the regimes considered here. When a merger becomes sufficiently strong that almost no coherent material survives, failure mode I dominates because the tracer population itself has disappeared. While usable cohorts remain, however, the largest mass errors are generally associated with failure mode II: the KLD-maximized spherical model does not reproduce the appropriate instantaneous BFE action-space structure. Failure mode III contributes an additional offset from the non-spherical spatial variation of the true force field, but this contribution remains mostly below the ten-percent level even in the 1:2 merger.

5. DISCUSSION AND CONCLUSION

5.1. Dependence on perturber orbital configuration

Figure 9 compares the standard, circular, and radial 1:5 merger configurations over their respective BFE-valid intervals. The radial configuration appears most favorable for the KLD fitting method over the periods analyzed here. Although the close radial encounter may produce a strong transient disturbance, the strongest forcing is plausibly concentrated into a relatively short interval. Outside this interval, coherent cohorts remain or reappear, and the fitted-minus-BFE action discrepancies and enclosed-mass biases are comparatively modest over much of the sampled evolution.

The circular configuration is less favorable in terms of sustained tracer availability. One possible interpretation is that prolonged merger-induced forcing continuously distorts the halo and stream population. The surviving action clouds generally have positive ΔD_{KL} and negative $\Delta \text{median}(J_r)$, indicating that the optimized spherical NFW model produces an action distribution that is more clustered and shifted toward lower radial actions relative to the BFE reference.

This orbital comparison should be regarded as suggestive rather than as a universal ranking of perturber trajectories, since the analysis is restricted to intervals in which BFE provides a sufficiently reliable reference. However, the comparison indicates that not only merger mass ratio, but also the duration and geometry of the encounter, affect the survival of coherent stream material and the form in which the action-space fitting failure appears.




Figure 9. Comparison between the KLD-fitted spherical-NFW model and the frozen-snapshot BFE reference potential in the $80 \leq r/\text{kpc} \leq 100$ radial range, restricted to snapshot intervals in which the BFE provides a reliable global reference. **Left column:** availability and stellar membership of the selected coherent cohorts. **Middle column:** signed fitted-minus-BFE discrepancies in the median radial action, $\Delta \text{median}(J_r)$, and in the radial-action clustering measure, $\Delta \text{KLD}(J_r)$. **Right column:** signed enclosed-mass bias of the fitted spherical model relative to the BFE reference. The radial configuration appears most favorable over the analyzed BFE-valid intervals, whereas the circular configuration appears least favorable for the continued survival of usable coherent cohorts.

5.2. Limitations of observable cohort selection and future prospects

The selection procedure introduced in Section 3 can effectively improve the KLD-maximization potential-constraining method by removing obviously unsuitable material. By retaining short, unfolded stream segments with smooth velocity profiles and locations away from orbital turning points, the method removes material that is most clearly phase-mixed or kinematically irregular. Nevertheless, it certainly does not provide a guarantee of unbiased potential recovery. The results of Section 4.3 demonstrate that apparently coherent cohorts may still yield biased fits when the optimized spherical potential does not reproduce the correct instantaneous action mapping or when the locally sampled gravitational field is not representative of the global spherical average.

These residual errors cannot generally be removed using the present selection criteria alone.

This limitation is especially relevant observationally because the true time-dependent potential and its corresponding instantaneous action distribution are not known. The comparisons from direct Dark Matter particle number-counting and from BFE potential models we present are therefore validation diagnostics only available in simulations, rather than quantities that can be measured directly for Galactic streams. Future work should investigate whether the onset of failure can instead be identified through observable or internally testable signatures, such as disagreement among independently fitted cohorts, sensitivity of the inferred potential to cohort removal, inconsistency between different radial ranges, or disagreement among multiple action-space diagnostics.

REFERENCES

- Arora, A., Sanderson, R. E., Panithanpaisal, N., et al. 2022, *ApJ*, 939, 2, doi: [10.3847/1538-4357/ac93fb](https://doi.org/10.3847/1538-4357/ac93fb)
- Besla, G., Kallivayalil, N., Hernquist, L., et al. 2007, *ApJ*, 668, 949, doi: [10.1086/521385](https://doi.org/10.1086/521385)
- Brooks, R. A. N., Garavito-Camargo, N., Johnston, K. V., et al. 2025, *ApJ*, 978, 79, doi: [10.3847/1538-4357/ad93a7](https://doi.org/10.3847/1538-4357/ad93a7)
- Brooks, R. A. N., Sanders, J. L., Lilleengen, S., Petersen, M. S., & Pontzen, A. 2024, *MNRAS*, 532, 2657, doi: [10.1093/mnras/stae1565](https://doi.org/10.1093/mnras/stae1565)
- Erkal, D., Li, T. S., Koposov, S. E., et al. 2018, *MNRAS*, 481, 3148, doi: [10.1093/mnras/sty2518](https://doi.org/10.1093/mnras/sty2518)
- Ferdosi, B. J., Buddelmeijer, H., Trager, S. C., Wilkinson, M. H. F., & Roerdink, J. B. T. M. 2011, *A&A*, 531, A114, doi: [10.1051/0004-6361/201116878](https://doi.org/10.1051/0004-6361/201116878)
- Kallivayalil, N., van der Marel, R. P., Besla, G., Anderson, J., & Alcock, C. 2013, *ApJ*, 764, 161, doi: [10.1088/0004-637X/764/2/161](https://doi.org/10.1088/0004-637X/764/2/161)
- Reino, S., Rossi, E. M., Sanderson, R. E., et al. 2021, *MNRAS*, 502, 4170, doi: [10.1093/mnras/stab304](https://doi.org/10.1093/mnras/stab304)
- Reino, S., Sanderson, R. E., Panithanpaisal, N., Rossi, E. M., & Kuijken, K. 2022, *MNRAS*, 509, 5365, doi: [10.1093/mnras/stab3176](https://doi.org/10.1093/mnras/stab3176)
- Sanderson, R. E., Helmi, A., & Hogg, D. W. 2015, *ApJ*, 801, 98, doi: [10.1088/0004-637X/801/2/98](https://doi.org/10.1088/0004-637X/801/2/98)
- Shipp, N., Li, T. S., Pace, A. B., et al. 2019, *ApJ*, 885, 3, doi: [10.3847/1538-4357/ab44bf](https://doi.org/10.3847/1538-4357/ab44bf)
- Vera-Ciro, C., & Helmi, A. 2013, *ApJL*, 773, L4, doi: [10.1088/2041-8205/773/1/L4](https://doi.org/10.1088/2041-8205/773/1/L4)
- Wu, X., Famaey, B., Gentile, G., Perets, H., & Zhao, H. 2008, *MNRAS*, 386, 2199, doi: [10.1111/j.1365-2966.2008.13198.x](https://doi.org/10.1111/j.1365-2966.2008.13198.x)