

Action-Angle Clustering of Stellar Streams in Merging Galaxies

MUQING WANG¹ AND TJITSKE STARKENBURG¹

¹*Northwestern University*

ABSTRACT

Stellar streams are sensitive tracers of the Galactic gravitational field, but action-clustering methods are formally grounded in time-independent integrable potentials, whereas the Milky Way is neither static nor isolated. We test the robustness of radial-action clustering using simulated globular-cluster streams evolved in live Milky Way-like halos undergoing mergers with different mass ratios and orbital configurations.

For each material cohort and simulation snapshot, we define a theoretical spherical targeted potential, $\Phi_{\text{tar},C}(t)$, whose force best matches the true acceleration field over the region occupied by that cohort. Freezing this potential supplies an integrable Hamiltonian and assigns a canonical instantaneous targeted action, $J_{\text{tar},C}(t)$, to every stellar phase-space point. Along the actual merger trajectories, the assigned action changes through two terms: the response to the residual force not represented by the frozen targeted potential and the remapping caused by evolution of the cohort-specific targeted-potential parameters. This framework motivates observable selection criteria that favor stream segments whose targeted-action clouds are likely to have undergone approximately semi-rigid translation rather than strong internal diffusion.

We select locally coherent cohorts using smooth along-stream velocity profiles, unfolded radial phase-space structure, limited spatial extent, sufficient radial span, and separation from radial turning points. For each selected cohort, we fit a spherical NFW potential by maximizing the one-dimensional Kullback–Leibler divergence (KLD) of its canonical fitted-model radial actions. At the distance-matched snapshot of the fiducial 1:10 merger, the selected cohorts recover a useful inner-halo enclosed-mass profile. Across the time evolution, stronger mergers generally remove usable coherent material earlier and produce larger and more rapidly varying mass biases, whereas the 1:100 merger remains comparatively stable.

We decompose the total fitted-to-true mass bias exactly into a fitted-to-targeted contribution and a targeted-to-true contribution. In the illustrated 1:10 inner-halo case, the former generally supplies the larger and more spatially extended structure. Broader targeted radial-action clouds permit a larger positive tail in the fitted-minus-targeted KLD and a wider range of fitted-to-targeted mass biases. Earlier snapshots preferentially occupy the large-KLD-excess and extreme-mass-bias envelopes, suggesting that an early but broadened stream can remain orderly enough to be artificially compressed by a mismatched trial potential, whereas later member-dependent disturbance is less readily remapped into a highly clustered action cloud. This interpretation concerns the accessibility of high clustering within the adopted trial family, not a universal late-time upper bound on KLD.

Trial-potential landscapes further reveal elongated regions with near-equivalent KLD scores and extended regions producing similar median radial actions. These degeneracies make the NFW parameters non-unique, but their mass consequence is directional: a KLD-preferred displacement along a nearly constant-enclosed-mass direction can leave the locally inferred mass nearly unchanged, whereas a displacement across enclosed-mass contours produces a mass bias. This motivates a post-fit mass-identifiability screen based on the range of enclosed masses spanned by the near-maximum KLD region. Finally, simulation-only comparisons between the fitted canonical radial actions and frozen-BFE radial-cycle integrations show time dependence related to the outer-halo mass bias. Observable cohort selection can therefore improve the robustness of action-clustering fits, but it cannot guarantee unique potential recovery or an unbiased global mass profile in a strongly perturbed halo.

Keywords: Galaxies (573) —

1. INTRODUCTION

Stellar streams form through the tidal disruption of globular clusters and dwarf galaxies. Because their stars

originate from dynamically cold progenitors and sub-

sequently follow neighboring trajectories in the Galactic gravitational field, the debris forms thin, coherent structures in phase space. Their small internal phase-space dispersion allows them to retain information about the mass profile, shape, and dynamical evolution of the Milky Way halo.

One way to extract this information is through the distribution of stream stars in action space. In a time-independent integrable Hamiltonian, canonical actions label invariant orbital tori and remain constant along each orbit. Stars stripped from a dynamically cold progenitor are therefore expected to occupy neighboring regions of action space when their phase-space coordinates are evaluated in an appropriate integrable potential. This motivates potential-fitting methods that search for the potential model in which stream actions are maximally clustered. [R. E. Sanderson et al. \(2015\)](#) formulated this approach using the Kullback–Leibler divergence (KLD), and subsequent studies have applied and tested it using mock and observed stream samples (e.g., [S. Reino et al. 2021, 2022](#)). More generally, the construction and numerical accuracy of action coordinates depend on the integrability or approximate separability of the adopted potential ([J. L. Sanders & J. Binney 2016](#)).

Maximum clustering is not, however, identical to recovery of the correct potential. An incorrect trial potential can shift or compress an action distribution, and a short stream segment with limited orbital-phase coverage may not distinguish between competing models. Potential recovery therefore depends not only on the intrinsic coldness of the tracer but also on its coverage in orbital phase, radius, energy, and angular momentum. The orbital-phase dependence found by [S. Reino et al. \(2021\)](#), together with the improved constraints obtained by combining dynamically distinct streams, illustrates this identifiability problem.

A live merger creates an additional and more fundamental complication. The true Galactic Hamiltonian is time dependent and generally non-integrable, so a global set of conserved canonical actions for the full system need not exist. We address this problem without treating the true potential as integrable. For each material cohort C and each instant, we instead define a cohort-specific spherical targeted potential, $\Phi_{\text{tar},C}(t)$, that best represents the true force over the region occupied by that cohort. Freezing this targeted potential defines an integrable action map for the instantaneous stellar phase-space points. As both the stars and their cohort-specific targeted potential evolve, the assigned targeted actions change through the residual force not reproduced by the targeted potential and through the change of the action map itself. Nearly common changes translate the action cloud without greatly altering its internal compactness, whereas member-dependent changes broaden or distort it.

This distinction motivates our observable cohort selection. We seek stream segments whose current morphology and kinematics indicate that their members probably experienced similar accumulated forcing and action-coordinate remapping. Smooth along-stream velocity profiles, unfolded radial phase-space branches, limited spatial extent, and separation from turning points cannot prove a semi-rigid targeted-action history, but they can reject material that is visibly folded, phase mixed, or kinematically irregular. The selection is therefore a physically motivated precondition for action clustering rather than a guarantee of an unbiased fit.

The Milky Way provides an important example of this non-equilibrium problem. Its halo is neither isolated nor exactly spherical, and it is perturbed by massive satellites and recent accretion events. In particular, the Large Magellanic Cloud (LMC) can generate a dynamical-friction wake and a large-scale response in the Milky Way halo ([C. Vera-Ciro & A. Helmi 2013](#); [N. Garavito-Camargo et al. 2019](#)). Its infall also induces reflex motion between the inner Galaxy and the outer halo, which can bias dynamical inferences that assume a stationary Galactocentric frame ([M. S. Petersen & J. Peñarrubia 2020](#)). Observed and simulated stellar streams exhibit track–velocity misalignments, orbital deflections, and other perturbations associated with the LMC and the deformation of the Milky Way halo (e.g., [D. Erkal et al. 2018](#); [N. Shipp et al. 2019](#); [R. A. N. Brooks et al. 2025](#)).

Recent studies have begun to connect these perturbations directly to action-based potential inference. [A. Arora et al. \(2022\)](#) showed that massive mergers can displace or disrupt stream structures in action space and that time-independent potential representations can lose orbital information on relatively short timescales. Using a time-evolving basis-function-expansion model of the Milky Way–LMC system, [R. A. N. Brooks et al. \(2024\)](#) found that radial-action clustering can retain useful local mass information for shorter-period streams without close LMC encounters, whereas longer-period and strongly perturbed streams may fail to recover the spherically averaged mass profile. These results motivate a distinction between canonical actions calculated in an integrable fitted model and radial-cycle diagnostics constructed from a more realistic but potentially non-integrable potential field.

In this work, we study these limitations using simulated globular-cluster streams evolved in live Milky Way-like halos undergoing mergers with different mass ratios and orbital configurations. For each selected cohort, we maximize the one-dimensional KLD of the canonical radial actions calculated in a trial spherical NFW potential. At the distance-matched snapshot of the fiducial 1:10 merger, the selected cohorts provide a useful recovery of the inner enclosed-mass profile. Repeating the procedure across snapshots and merger strengths then

shows when coherent tracers disappear and when surviving cohorts cease to yield stable mass estimates.

We separate these outcomes into three logically distinct levels. First, a perturbed snapshot may contain too few observable coherent cohorts to constrain the potential. Second, even when such cohorts survive, KLD maximization may select a spherical potential different from the cohort-wise targeted potential. We diagnose this fitted-to-targeted failure through artificial over-clustering and through degeneracies of the KLD score and median action across the trial-parameter plane. Third, even exact recovery of the targeted spherical potential need not reproduce the global spherically averaged mass profile. An exact signed mass-bias decomposition separates the second and third levels.

Our parameter-plane analysis also distinguishes potential-parameter identifiability from mass identifiability. Different NFW parameters may occupy an extended near-maximum KLD ridge while predicting similar enclosed masses over the cohort's radial range. Conversely, a KLD-preferred displacement that crosses enclosed-mass contours produces a mass bias. This directional dependence motivates a practical post-fit screen: map the near-maximum KLD region into enclosed-mass space and favor cohorts for which it corresponds to a narrow mass range.

Finally, we use simulation-only information to compare the fitted canonical radial actions with radial-cycle integrations in the frozen full BFE potential and to study the dependence on perturber orbital configuration. These diagnostics do not enter the observable fitting method; they test how the fitted spherical orbital mapping differs from the full simulated field.

Section 2 develops the canonical-action, cohort-wise targeted-potential, and targeted-action-evolution framework. Section 3 presents the simulations, observable selection criteria, and baseline performance. Section 4 decomposes the post-selection mass error and analyzes over-clustering and degeneracy. Section 5 presents diagnostics that require knowledge of the simulated full potential. Section 6 discusses limitations, observational use, and future extensions.

2. THEORETICAL FRAMEWORK

2.1. Canonical Actions and Potential-dependent Action Clustering

For a Hamiltonian system with n degrees of freedom, the canonical action variables are well defined for regular motion confined to invariant tori. In particular, Liouville integrability requires n functionally independent integrals of motion F_i that are mutually in involution,

$$\{F_i, F_j\} = 0 \quad \text{for all } i, j. \quad (1)$$

When the corresponding common level set is regular, compact, and connected, it forms an invariant n -torus. The canonical actions are then defined by integrals over

the n independent fundamental cycles γ_i of this torus:

$$J_i = \frac{1}{2\pi} \oint_{\gamma_i} \mathbf{p} \cdot d\mathbf{q}. \quad (2)$$

In separable canonical coordinates, this expression reduces to

$$J_i = \frac{1}{2\pi} \oint_{\gamma_i} p_i dq_i. \quad (3)$$

Here, q_i and p_i are canonically conjugate variables. The actions label the invariant torus and remain constant under the integrable Hamiltonian, while the conjugate angles evolve along it.

A static spherically symmetric potential is integrable. Away from degenerate orbits, the energy H , total angular momentum squared L^2 , and azimuthal angular momentum L_z provide three functionally independent integrals in involution. Consequently, three canonical actions (J_r, J_θ, J_ϕ) can be defined. Throughout the following expressions, momenta and energies are taken per unit stellar mass. The energy is

$$\begin{aligned} E &= \frac{1}{2} \left(p_r^2 + \frac{p_\theta^2}{r^2} + \frac{p_\phi^2}{r^2 \sin^2 \theta} \right) + \Phi(r) \\ &= \frac{1}{2} \left(p_r^2 + \frac{L^2}{r^2} \right) + \Phi(r), \end{aligned} \quad (4)$$

where

$$L^2 = p_\theta^2 + \frac{p_\phi^2}{\sin^2 \theta}. \quad (5)$$

Since the azimuthal angle ϕ is cyclic in a spherical potential, its conjugate momentum is conserved:

$$p_\phi = L_z. \quad (6)$$

So that the azimuthal action is:

$$J_\phi = \frac{1}{2\pi} \oint_{\gamma_\phi} p_\phi \cdot d\phi = p_\phi = L_z \quad (7)$$

Using the expression of angular momentum shown in Equ.4, the polar action can be expressed as:

$$J_\theta = \frac{1}{2\pi} \oint_{\gamma_\theta} p_\theta \cdot d\theta = \frac{1}{\pi} \int_{\theta_{min}}^{\theta_{max}} \sqrt{L^2 - \frac{L_z^2}{\sin^2 \theta}} \cdot d\theta \quad (8)$$

Using $\sin(\theta_{min}) = \sin(\theta_{max}) = \frac{|L_z|}{L}$, the integral gives:

$$J_\theta = L - |L_z| \quad (9)$$

Finally, using Equ.4, the radial action can also be acquired:

$$J_r = \frac{1}{\pi} \int_{r_{min}}^{r_{max}} \sqrt{2[E - V(r)] - \frac{L^2}{r^2}} \cdot dr \quad (10)$$

where r_{min} and r_{max} are the radial turning points in the adopted spherical potential. For the fitted static spherical halo models, we calculate these canonical actions

using **Agama**. The code determines the allowed radial interval and evaluates the corresponding canonical actions.

Stars stripped from a dynamically cold progenitor begin on nearby orbits and therefore occupy a relatively compact region in action space when evaluated in an appropriate integrable potential. This compactness contains information about the potential. Stars from the same stream generally sample different orbital phases. *An appropriate potential maps these different phase-space points onto nearby invariant tori*, whereas an incorrect trial potential generally assigns phase-dependent action errors, broadening or distorting the inferred action cloud. This is the physical basis of action-clustering potential inference (R. E. Sanderson et al. 2015; S. Reino et al. 2021, 2022).

This statement is not an identity: a wrong potential may translate an action cloud without greatly broadening it, and a short stream segment may provide insufficient phase coverage to distinguish trial models. Potential recovery therefore requires both a dynamically cold tracer and sufficient diversity in orbital phase, radius, energy, or angular momentum. Combining multiple streams on different orbits can reduce individual-stream degeneracies and biases (R. E. Sanderson et al. 2015; S. Reino et al. 2021).

2.2. 1D Kullback-Leibler Divergence

We quantify radial-action clustering using the Kullback–Leibler divergence (KLD). For a radial-action probability density $p(J_r)$ and a reference density $q(J_r)$,

$$D_{\text{KL}}[p||q] = \int p(J_r) \ln \left[\frac{p(J_r)}{q(J_r)} \right] dJ_r. \quad (11)$$

where $p(J_r)$ is the probability density of the radial action distribution at the points J_r , and $q(J_r)$ is the background distribution.

Also, following the work by R. E. Sanderson et al. (2015), I use the Modified Breiman Density Estimator to calculate the density profile of the action points. According to B. J. Ferdosi et al. (2011), a fine grid will be built across the action space, and the density value at the center of each fine grid $J_{r,k}$ will be evaluated using the following formula:

$$p(J_{r,k}) = \frac{1}{N} \sum_{i=1}^N (\sigma \lambda_i)^{-d} K_E \left(\frac{J_{r,k} - J_{r,i}}{\sigma_r \lambda_i} \right) \quad (12)$$

Here, i is the star index, N stands for the total number of stars, and d stands for the dimensionality of the action space. This estimator uses the Epanechnikov Kernel K_E instead of a Gaussian Kernel:

$$K_E(t) = \begin{cases} \frac{d+2}{2V_d} (1-t^2), & \text{if } t^2 < 1, \\ 0, & \text{otherwise.} \end{cases} \quad (13)$$

, where V_d is the volume of the unit sphere in d -dimensional space. We can easily see the kernel contribution decreases with normalized distance from the sample point to the grid center, which is as expected. Following the analysis in B. J. Ferdosi et al. (2011), the window sizes σ_j in each direction in the action space is derived from coordinate percentiles in each dimension in the action space:

$$\sigma_r = \min \left\{ \frac{P_{80}(J_{r,i}) - P_{20}(J_{r,i})}{\ln N} \right\} \quad (14)$$

, where $P_n(J_{r,i})$ is the n th percentile of the radial action values. In addition, define a pilot density $p_{\text{pilot}}(J_{r,i})$ to describe the degree of clustering for the data point by using Equ. 12, with $\lambda_i = 1$. And the adaptive bandwidth λ_i of this estimator is designed to be inversely proportional to the $p_{\text{pilot}}(J_{r,i})$:

$$\lambda_i = \left(\frac{p_{\text{pilot}}(J_{r,i})}{g} \right)^{-\alpha} \quad (15)$$

, where g is the geometric mean of the pilot density:

$$g = \exp \left[\frac{1}{N} \sum_{i=1}^N \ln p_{\text{pilot}}(J_{r,i}) \right] \quad (16)$$

and $\alpha = \frac{1}{d}$ is the adaptive-bandwidth exponent. This adaptive bandwidth shrinks kernels in dense regions and widens them in sparse regions, thus increasing the density value where the action points cluster tightly.

Also, for those grids with overly small $p(J_r)$ values, a floor value (like 10^{-9}) needs to be assigned to prevent error. Luckily, the function “`modified_breiman_density`” in the “`MBEtree`” package can realize the above calculation for $p(J_r)$.

Following the action-clustering framework of R. E. Sanderson et al. (2015) and its applications by S. Reino et al. (2021, 2022), trial gravitational potentials are compared through the degree of clustering of the actions calculated from the same observed phase-space data. In this work, the reference distribution $q(J_r)$ is uniform over one fixed and sufficiently wide action interval $[J_{r,\min}, J_{r,\max}]$ that is shared by every trial potential,

$$q(J_r) = \frac{1}{J_{r,\max} - J_{r,\min}}. \quad (17)$$

The interval $[J_{r,\min}, J_{r,\max}]$ is kept the same for all trial potentials. In practice, we choose this universal range to be sufficiently wide to contain the radial actions produced over the full set of potentials considered, including deliberately extreme choices of the potential parameters. This ensures that changes in the KLD are not driven by changing the support of the background distribution from one potential to another.

Thus, the KLD measures how strongly the radial-action distribution departs from a uniform distribution

over the fixed reference range. Since the reference interval is held fixed for all trial potentials, maximizing D_{KL}^I is equivalent to finding the potential in which the stream is most concentrated in J_r , or equivalently, has the lowest radial-action entropy. A mismatched potential generally broadens the inferred J_r distribution of a stream, thereby reducing the KLD. The KLD test used here is based on the expectation that the most accurate potential produces the sharpest radial-action clustering relative to the same uniform background. The success of this method has been demonstrated in static test potentials and in applications combining multiple stellar streams (R. E. Sanderson et al. 2015; S. Reino et al. 2021, 2022).

2.3. Cohort-wise Snapshot Targeted Potentials

A live merger differs fundamentally from the static integrable setting. Its true Hamiltonian,

$$H(\mathbf{z}, t) = \frac{1}{2}|\mathbf{p}|^2 + \Phi(\mathbf{x}, t), \quad \mathbf{z} \equiv (\mathbf{x}, \mathbf{p}), \quad (18)$$

is time dependent and generally non-integrable. The true trajectories therefore need not lie on invariant three-tori, and the full merger potential does not generally possess a global set of conserved canonical actions.

The fitting problem nevertheless requires an integrable model that can assign actions to the observed phase-space points. We therefore define a theoretical spherical targeted potential separately for every material cohort C and every instant t ,

$$\Phi_{\text{tar},C}(r, t) = \Phi_{\text{sph}}[r; \boldsymbol{\lambda}_{\text{tar},C}(t)]. \quad (19)$$

Here C denotes a fixed set of stellar members when its history is considered. Different cohorts generally occupy different regions of the force field and therefore possess different time-indexed sequences of targeted parameters.

In the simulations, a precise theoretical definition is obtained by minimizing the force mismatch over the cohort,

$$\boldsymbol{\lambda}_{\text{tar},C}(t) = \arg \min_{\boldsymbol{\lambda}} \frac{1}{N_C} \sum_{a \in C} w_a \|\mathbf{g}[\mathbf{x}_a(t), t] - \mathbf{g}_{\text{sph}}[r_a(t); \boldsymbol{\lambda}]\|^2, \quad (20)$$

where $\mathbf{g} = -\nabla\Phi$, $\mathbf{g}_{\text{sph}} = -\nabla\Phi_{\text{sph}}$, and w_a is the adopted stellar weight. Because every candidate is spherical, only its radial force can be adjusted; tangential residual forces remain part of the mismatch. Equation 20 defines the spherical model that best represents the local force sampled by cohort C . It does not make the true local potential spherical or integrable.

At a snapshot t_n , freezing any candidate value of $\boldsymbol{\lambda}$ gives a static integrable Hamiltonian,

$$H_{\text{sph}}(\mathbf{z}; \boldsymbol{\lambda}) = \frac{1}{2}|\mathbf{p}|^2 + \Phi_{\text{sph}}(r; \boldsymbol{\lambda}), \quad (21)$$

which assigns a regular model orbit and canonical actions $\mathbf{J}(\mathbf{z}; \boldsymbol{\lambda})$ to each stellar phase-space point. The

snapshot-wise targeted action of member a is consequently

$$J_{\text{tar},C,i}^{(a)}(t) \equiv J_i[\mathbf{z}_a(t); \boldsymbol{\lambda}_{\text{tar},C}(t)]. \quad (22)$$

It is a canonical action of the instantaneously frozen targeted Hamiltonian, but it is not generally conserved along the star's actual trajectory.

For a given material cohort, Equations 20 and 22 define the snapshot sequence

$$\Phi_{\text{tar},C}(t_1) \rightarrow \Phi_{\text{tar},C}(t_2) \rightarrow \cdots \rightarrow \Phi_{\text{tar},C}(t_n). \quad (23)$$

This sequence is a chain of cohort-wise approximations, not the physical evolution of an integrable true potential. When the minimum in Equation 20 follows an isolated smooth branch, $\boldsymbol{\lambda}_{\text{tar},C}(t)$ may be differentiated locally. If the optimum changes discontinuously between branches, the finite-difference construction below remains valid, whereas its continuous-time form applies only piecewise.

2.4. Hamiltonian Evolution of the Assigned Targeted Actions

For brevity, consider one member of cohort C and suppress the labels a and C temporarily. The assigned action is the composite function

$$J_{\text{tar},i}(t) = J_i[\mathbf{z}(t); \boldsymbol{\lambda}_{\text{tar}}(t)]. \quad (24)$$

Its total derivative along the actual stellar trajectory is

$$\frac{dJ_{\text{tar},i}}{dt} = \left. \frac{\partial J_i}{\partial \mathbf{z}} \right|_{\boldsymbol{\lambda}_{\text{tar}}} \cdot \dot{\mathbf{z}} + \dot{\lambda}_{\text{tar}}^\alpha \left. \frac{\partial J_i}{\partial \lambda_{\text{tar}}^\alpha} \right|_{\mathbf{z}}. \quad (25)$$

The first term is the change caused by the star's actual phase-space motion while the instantaneous targeted action map is held fixed. The second is the change caused by updating that map at a fixed physical phase-space point. Repeated targeted-potential parameter indices α are summed.

The first term follows directly from Hamilton's equations. Holding $\boldsymbol{\lambda}_{\text{tar}}$ fixed,

$$\begin{aligned} \left. \frac{dJ_i}{dt} \right|_{\boldsymbol{\lambda}_{\text{tar}}} &= \sum_k \left(\frac{\partial J_i}{\partial x_k} \dot{x}_k + \frac{\partial J_i}{\partial p_k} \dot{p}_k \right) \\ &= \sum_k \left(\frac{\partial J_i}{\partial x_k} \frac{\partial H}{\partial p_k} - \frac{\partial J_i}{\partial p_k} \frac{\partial H}{\partial x_k} \right) = \{J_i, H\}, \end{aligned} \quad (26)$$

where $\dot{x}_k = \partial H / \partial p_k$ and $\dot{p}_k = -\partial H / \partial x_k$ are generated by the true Hamiltonian.

At that instant, decompose the true Hamiltonian into the frozen targeted Hamiltonian and a residual,

$$H(\mathbf{z}, t) = H_{\text{tar},C}[\mathbf{z}; \boldsymbol{\lambda}_{\text{tar},C}(t)] + \delta H_C(\mathbf{z}, t), \quad (27)$$

with

$$\delta H_C = \Phi(\mathbf{x}, t) - \Phi_{\text{tar},C}[r; \boldsymbol{\lambda}_{\text{tar},C}(t)] \equiv \delta \Phi_C. \quad (28)$$

In the action–angle coordinates of the frozen targeted Hamiltonian, $H_{\text{tar},C} = H_{\text{tar},C}(\mathbf{J}; \boldsymbol{\lambda}_{\text{tar},C})$ has no angle dependence. Therefore,

$$\begin{aligned} \{J_i, H\} &= \{J_i, H_{\text{tar},C}\} + \{J_i, \delta H_C\} \\ &= -\frac{\partial H_{\text{tar},C}}{\partial \theta_{\text{tar},C,i}} - \frac{\partial \delta H_C}{\partial \theta_{\text{tar},C,i}} = -\frac{\partial \delta \Phi_C}{\partial \theta_{\text{tar},C,i}}. \end{aligned} \quad (29)$$

Defining the residual force

$$\delta \mathbf{F}_C \equiv \mathbf{g} - \mathbf{g}_{\text{tar},C} = -\nabla \delta \Phi_C, \quad (30)$$

where $\mathbf{g}_{\text{tar},C} \equiv -\nabla \Phi_{\text{tar},C}$, the chain rule gives

$$\begin{aligned} -\frac{\partial \delta \Phi_C}{\partial \theta_{\text{tar},C,i}} &= -\nabla \delta \Phi_C \cdot \frac{\partial \mathbf{x}}{\partial \theta_{\text{tar},C,i}} \\ &= \delta \mathbf{F}_C \cdot \frac{\partial \mathbf{x}}{\partial \theta_{\text{tar},C,i}}. \end{aligned} \quad (31)$$

Restoring the cohort and member labels in Equation 25 yields

$$\boxed{\frac{dJ_{\text{tar},C,i}^{(a)}}{dt} = \delta \mathbf{F}_C^{(a)} \cdot \frac{\partial \mathbf{x}}{\partial \theta_{\text{tar},C,i}} + \dot{\lambda}_{\text{tar},C}^\alpha \frac{\partial J_{\text{tar},C,i}^{(a)}}{\partial \lambda_{\text{tar},C}^\alpha} \Big|_{\mathbf{z}_a}}. \quad (32)$$

The connection to neighboring simulation snapshots is exact. Adding and subtracting the action assigned to the new phase-space point by the old targeted potential gives

$$\begin{aligned} J_i(\mathbf{z}_{a,n+1}; \boldsymbol{\lambda}_{\text{tar},C,n+1}) - J_i(\mathbf{z}_{a,n}; \boldsymbol{\lambda}_{\text{tar},C,n}) \\ = \left[J_i(\mathbf{z}_{a,n+1}; \boldsymbol{\lambda}_{\text{tar},C,n}) - J_i(\mathbf{z}_{a,n}; \boldsymbol{\lambda}_{\text{tar},C,n}) \right] \\ + \left[J_i(\mathbf{z}_{a,n+1}; \boldsymbol{\lambda}_{\text{tar},C,n+1}) - J_i(\mathbf{z}_{a,n+1}; \boldsymbol{\lambda}_{\text{tar},C,n}) \right]. \end{aligned} \quad (33)$$

The first bracket evaluates the star’s two actual phase-space points with the same frozen targeted action map; it does not evolve the star under the targeted Hamiltonian. Dividing by Δt and taking the smooth limit gives

$$\frac{J_i(\mathbf{z}_{a,n+1}; \boldsymbol{\lambda}_{\text{tar},C,n}) - J_i(\mathbf{z}_{a,n}; \boldsymbol{\lambda}_{\text{tar},C,n})}{\Delta t} \rightarrow \delta \mathbf{F}_C^{(a)} \cdot \frac{\partial \mathbf{x}}{\partial \theta_{\text{tar},C,i}}. \quad (34)$$

At the common endpoint $\mathbf{z}_{a,n+1}$, a first-order Taylor expansion of the second bracket gives

$$\frac{J_i(\mathbf{z}_{a,n+1}; \boldsymbol{\lambda}_{\text{tar},C,n+1}) - J_i(\mathbf{z}_{a,n+1}; \boldsymbol{\lambda}_{\text{tar},C,n})}{\Delta t} \rightarrow \dot{\lambda}_{\text{tar},C}^\alpha \frac{\partial J_i}{\partial \lambda_{\text{tar},C}^\alpha} \Big|_{\mathbf{z}_a} \quad (35)$$

Equation 32 is therefore the continuous-time limit of comparing the actions assigned along the cohort’s sequence of snapshot-wise spherical targeted potentials.

2.5. From Targeted-action Evolution to Cohort Selection

The action cloud remains useful for clustering only if the evolution in Equation 32 is sufficiently similar among the members. For two members a and b ,

$$\begin{aligned} \frac{d}{dt} \left(J_{\text{tar},C,i}^{(a)} - J_{\text{tar},C,i}^{(b)} \right) &= \Delta_{ab} \left[\delta \mathbf{F}_C \cdot \frac{\partial \mathbf{x}}{\partial \theta_{\text{tar},C,i}} \right] \\ &\quad + \dot{\lambda}_{\text{tar},C}^\alpha \Delta_{ab} \left[\frac{\partial J_{\text{tar},C,i}}{\partial \lambda_{\text{tar},C}^\alpha} \Big|_{\mathbf{z}} \right], \end{aligned} \quad (36)$$

where $\Delta_{ab}X \equiv X^{(a)} - X^{(b)}$. Integrating between an earlier time t_1 and the observed time t_0 gives

$$\Delta_{ab}J_{\text{tar},C,i}(t_0) = \Delta_{ab}J_{\text{tar},C,i}(t_1) + \int_{t_1}^{t_0} \Delta_{ab}\dot{J}_{\text{tar},C,i}(t) dt. \quad (37)$$

If the total action change is nearly common across the cohort, the cloud may translate through action space while retaining its internal compactness. If the changes are strongly member dependent, the cloud broadens or distorts even under the correct instantaneous targeted potential. KLD maximization can then prefer a different trial potential that artificially compresses the disturbed distribution.

The force history and the sequence $\Phi_{\text{tar},C}(t)$ are not directly observable. We therefore use present-day phase-space structure as evidence, rather than proof, of a semi-rigid past evolution. A cohort that is smooth in velocity, occupies one unfolded phase-space branch, spans a restricted region, and lies away from a radial turning point is more likely to have experienced similar residual-force projections and similar action-map changes among its members. Passing the selection at t_0 is consequently a strong indicator that the instantaneous targeted-action cloud remains compact enough to support clustering inference.

The selected cohort population may differ from one snapshot to another without invalidating this construction. Every cohort selected at a given snapshot has its own material history and its own hypothetical sequence of cohort-wise targeted potentials. The time-dependent population plots later in the paper are repeated cross-sectional applications of the same criterion; they do not assert that one fixed cohort survives the selection at every snapshot.

3. SIMULATIONS, COHORT SELECTION, AND BASELINE PERFORMANCE

3.1. Merger Simulations and BFE Representation

We use the stellar-stream simulations from *Dancing Streams in Merging Halos* (Weerasooriya et al. 2025), a controlled suite designed to study how globular-cluster streams respond to a Milky Way–LMC-like merger. The

simulations first generate 1024 mock globular-cluster streams in a static Milky Way-like Hernquist halo using the Fardal et al. (2015) mock-stream distribution function as implemented in *Gala*. Each stream contains 1000 stellar particles. The grid spans progenitor masses of 10^4 and $10^6 M_\odot$, eight apocentric radii from 30 to 100 kpc, eight circularities from $\eta = 0.5$ to 1.0, and eight stream ages from 0.5 to 6 Gyr.

The same stream population is then inserted into live N -body halos and evolved with *Gadget4*. The host halo has mass $M_{\text{halo}} = 1.57 \times 10^{12} M_\odot$ and Hernquist scale radius $a = 40.85$ kpc. In addition to the published isolated and 1:5 simulations, we use related data products spanning perturber-to-host mass ratios 1:2, 1:5, 1:10, and 1:100, together with standard, more circular, and more radial 1:5 orbital configurations. The simulations contain 241 snapshots over 6 Gyr; the principal KLD analysis uses snapshots 000–100, corresponding to approximately 0–2.56 Gyr.

The time-dependent halo field is represented by a basis-function expansion (BFE),

$$\rho(\mathbf{x}, t) \simeq \sum_{n, \ell, m} a_{n\ell m}(t) u_{n\ell}(r) Y_{\ell m}(\theta, \phi), \quad (38)$$

with the corresponding potential obtained from Poisson’s equation. The adopted coefficients were supplied by David Amir Gonzalez and Robyn Sanderson and accurately represent the later-stage potential over the BFE-valid intervals used below. The BFE provides the simulated full force field required to define the cohort-wise targeted potentials and the simulation-only diagnostics; it is not used as an observable input to the KLD fit.

3.2. Observable Cohort-selection Criteria

Neither the residual-force history nor the targeted-potential chain is available in an observation. We therefore implement the theoretical motivation of Section 2.5 through present-day phase-space criteria applied before the KLD fit.

First, we require a smooth velocity profile along the stream. For an along-stream coordinate s , we fit

$$v_{\text{fit}}(s) = As^2 + Bs + C \quad (39)$$

and retain only cohorts satisfying

$$\Delta_v \equiv \frac{\sqrt{\langle [v_{\parallel}(s) - v_{\text{fit}}(s)]^2 \rangle}}{\text{std}[v_{\parallel}(s)]} < 0.1. \quad (40)$$

A smooth profile is consistent with a small member-dependent accumulated velocity response, although it does not uniquely determine the past force history.

Second, we require an unfolded, single-valued radial phase-space branch. Operationally, $v_{\text{fit}}(s)$ must remain

strictly monotonic over the selected interval. This rejects folded structures in which nearby positions correspond to different orbital branches and therefore potentially different targeted-orbit tangents and action sensitivities.

Third, the cohort must have a limited spatial extent, so that its members sample a restricted region of the force field and can be represented by one cohort-common targeted potential. An extremely narrow radial segment, however, has little leverage to distinguish trial force profiles. We therefore additionally require

$$\Delta r_{\text{cohort}} \equiv \max_a(r_a) - \min_a(r_a) \geq 2 \text{ kpc}. \quad (41)$$

The upper extent restriction is a coherence condition, whereas the lower radial-span restriction is an identifiability condition. All present-day candidate cohorts considered here already satisfy the latter.

Fourth, we exclude cohorts close to radial turning points by imposing a threshold on $|v_r|$. Near pericentre or apocentre, the mapping between local phase-space coordinates and radial phase becomes poorly conditioned, and incoming and outgoing branches can have substantially different responses despite occupying similar radii.

Finally, every retained cohort must contain at least 100 stars. The four physical criteria are thus smooth along-stream velocity, an unfolded phase-space branch, limited spatial extent with sufficient radial span, and separation from radial turning points. Figure 1 illustrates that this procedure is deliberately restrictive: it retains short, nearly single-valued segments rather than merely removing isolated outliers. Passing these criteria favors, but does not guarantee, a compact targeted-action cloud and reliable potential recovery.

3.3. Implementation and Distance-matched Present-day Recovery

We first implement the selection in the fiducial 1:10 merger. Because the LMC is currently at a Galactocentric distance of approximately 49.5 kpc (X. Wu et al. 2008) and is likely on its first infall or first passage (G. Besla et al. 2007; N. Kallivayalil et al. 2013), we identify snapshot 55 ($t = 1.406$ Gyr) as the distance-matched present-day analogue. At this snapshot the simulated host–perturber separation first crosses the observed value.




Figure 1. Illustration of the cohort-selection procedure in radial phase space for the 1:10 merger at snapshot 030 in the $30 \leq r/\text{kpc} \leq 40$ radial range. The vertical dashed lines mark the boundaries of the selected radial interval. **Left:** four parent streams together with six comparison streams before cohort selection, including extended, folded, and kinematically irregular phase-space structures. **Right:** the locally coherent segments retained from the four parent streams after requiring a smooth along-stream velocity profile, an unfolded monotonic radial phase-space branch, limited spatial extent, and separation from radial turning points. The listed values of N give the stellar membership of each retained cohort. These criteria favor segments with smaller member-to-member differential forcing and targeted-coordinate remapping, but do not by themselves guarantee an unbiased potential fit.




Figure 2. Host-perturber separation as a function of snapshot number for the fiducial Milky Way-LMC-like merger. The horizontal dashed line marks the target separation of 49.5 kpc, and the vertical dashed line identifies its first inbound crossing at snapshot 055 ($t = 1.406$ Gyr). We adopt this snapshot as the distance-matched present-day analogue used for the single-snapshot cohort-selection and enclosed-mass analysis.

For every selected cohort, we maximize the radial-action KLD over a spherical NFW family and obtain an independent enclosed-mass profile. The initial halo is spherical, but the live merger field later becomes non-spherical; the fitted NFW potential is therefore an action-computable approximation, not a claim that the true halo remains spherical. We retain cohorts containing at least 100 stars in successive 10-kpc-wide radial ranges. No cohort passes the adopted criteria inside 10 kpc, where the simulated perturbations are especially strong.

The spherical enclosed mass implied by an arbitrary gravitational potential can be obtained from the radial gravitational acceleration on a spherical surface. The gravitational acceleration is

$$\mathbf{g}(\mathbf{x}, t) = -\nabla\Phi(\mathbf{x}, t), \quad (42)$$

while Poisson's equation gives

$$\nabla^2\Phi(\mathbf{x}, t) = 4\pi G\rho(\mathbf{x}, t). \quad (43)$$

Combining Equations 42 and 43, we obtain

$$\nabla \cdot \mathbf{g}(\mathbf{x}, t) = -4\pi G\rho(\mathbf{x}, t). \quad (44)$$

We integrate Equation 44 over the volume V_r enclosed by a sphere of radius r :

$$\begin{aligned} \int_{V_r} \nabla \cdot \mathbf{g} d^3\mathbf{x} &= -4\pi G \int_{V_r} \rho(\mathbf{x}, t) d^3\mathbf{x} \\ &= -4\pi G M_{\text{sph}}(< r, t). \end{aligned} \quad (45)$$

Applying the divergence theorem gives

$$\oint_{S_r} \mathbf{g} \cdot d\mathbf{S} = -4\pi G M_{\text{sph}}(< r, t), \quad (46)$$

where S_r is the spherical surface bounding V_r .

On a sphere of radius r , the outward surface element is

$$d\mathbf{S} = r^2 \hat{\mathbf{r}} d\Omega, \quad (47)$$

and the radial component of the gravitational acceleration is

$$g_r(r, \Omega, t) \equiv \mathbf{g}(r, \Omega, t) \cdot \hat{\mathbf{r}}. \quad (48)$$

Also, we can define the angular average of the radial acceleration at fixed radius as

$$\langle g_r(r, \Omega, t) \rangle_\Omega \equiv \frac{1}{4\pi} \int_{4\pi} g_r(r, \Omega, t) d\Omega. \quad (49)$$

Substituting Equations 47, 48 and 49 into Equation 46 yields the spherically averaged enclosed mass:

$$M_{\text{sph}}(< r, t) = -\frac{r^2}{G} \langle g_r(r, \Omega, t) \rangle_\Omega. \quad (50)$$

The above equation does not require the gravitational potential to be spherical. It follows from the total gravitational flux through a spherical surface and therefore returns the mass enclosed by that surface implied by the potential model.

The angular average in Equation 50 is evaluated numerically using $N_\Omega = 4096$ nearly uniformly distributed Fibonacci-sphere directions. Thus,

$$\langle g_r(r, \Omega, t) \rangle_\Omega \simeq \frac{1}{N_\Omega} \sum_{a=1}^{N_\Omega} g_r(r, \Omega_a, t), \quad (51)$$

and the enclosed mass is estimated as

$$M_{\text{sph}}(< r, t) \simeq -\frac{r^2}{G N_\Omega} \sum_{a=1}^{N_\Omega} g_r(r, \Omega_a, t). \quad (52)$$

The Fibonacci construction provides nearly equal-area sampling over the sphere and is sufficiently dense relative to the angular resolution of the NFW spherical model.

We combine the independent cohort estimates through a member-number-weighted enclosed-mass profile. In the left panel of Figure 3, we plot the enclosed-mass profile from each cohort only over the radial interval occupied by that cohort. We divide the full radial range

into 0.1-kpc-wide bins and, at each bin r_k , average the fitted enclosed masses of the contributing cohorts:

$$M_{\text{fit}}(< r_k, t) = \frac{\sum_{i \in \mathcal{I}_k} N_i M_{\text{sph}, i}(< r_k, t)}{\sum_{i \in \mathcal{I}_k} N_i}, \quad (53)$$

where $\mathcal{I}_k$ denotes the selected cohorts whose fitted profiles are evaluated at r_k , and N_i is the number of selected stars in cohort i . The right panel of Figure 3 shows the weight-averaged enclosed mass value at each radial bin from the galactic center outward.

The resulting profile shows that, despite the spherical model simplification, the selected streams recover the enclosed mass reasonably well over $10 \lesssim r/\text{kpc} \lesssim 30$. The discrepancy generally increases farther outward. At snapshot 55, the perturber lies about 49.5 kpc from the host center, and coherent outer-halo segments preferentially survive in less disturbed regions farther from it. A locally constrained spherical model in those regions need not represent the global spherical average and can consequently underestimate the total enclosed mass. At this merger stage, the selected-cohort KLD method is therefore most informative over the inner radial range in which the individual cohort constraints overlap.

3.4. Time-dependent Inner-halo Performance

The distance-matched snapshot demonstrates that the selected cohorts can recover a useful enclosed-mass profile at a particular stage of the merger. However, this single-snapshot result does not establish how long the method remains reliable as the halo and stream population evolve. We therefore repeat the cohort selection and spherical-NFW fitting at every snapshot from 0 to 100, corresponding to $0 \leq t \lesssim 2.56$ Gyr.

We consider cohorts occupying the radial interval $30 \leq r/\text{kpc} \leq 40$, matching Figure 4. For this inner-halo analysis, the true enclosed mass is measured directly from the host and perturber dark-matter particles, rather than from the BFE potential.

To characterize the surviving stream sample independently of the tested potential, we define the six-dimensional phase-space compactness of cohort i as

$$C_i = \left[\frac{1}{N_i} \sum_{j=1}^{N_i} \|\mathbf{z}_{ij} - \bar{\mathbf{z}}_i\|^2 \right]^{1/2}, \quad (54)$$

where j is the star index, N_i is the number of selected stars in cohort i , and

$$\mathbf{z} = \left(\frac{x}{10 \text{ kpc}}, \frac{y}{10 \text{ kpc}}, \frac{z}{10 \text{ kpc}}, \frac{v_x}{100 \text{ km s}^{-1}}, \frac{v_y}{100 \text{ km s}^{-1}}, \frac{v_z}{100 \text{ km s}^{-1}} \right). \quad (55)$$

The selected-cohort-weighted compactness is

$$\langle C \rangle_w = \frac{\sum_i N_i C_i}{\sum_i N_i}. \quad (56)$$




Figure 3. Enclosed-mass recovery from the selected coherent cohorts at snapshot 055, corresponding to the distance-matched present-day configuration of the fiducial 1:10 MW–LMC-like merger. **(a)** KLD-maximized spherical NFW enclosed-mass profiles inferred independently from the individual selected cohorts. Each grey curve is shown only over the radial interval sampled by its corresponding cohort, while the black curve shows the true dark-matter enclosed-mass profile measured directly from the snapshot particles relative to the host center. **(b)** Member-number-weighted mean of the independently fitted cohort mass profiles at each radius (red), compared with the same true enclosed-mass profile (black). Only cohorts whose sampled radial interval includes the given radius contribute to the weighted mean. Thus, the red curve is a radial combination of cohort-specific fits, rather than the result of a single joint fit to all selected stars. The recovered profile agrees most closely with the true profile at smaller radii, while the discrepancy generally increases farther outward. This behavior is consistent with the increasing importance of merger-induced departures from a spherical equilibrium model, together with the decreasing availability and radial coverage of useful coherent cohorts at large radii.

At each radial grid point r_k within the 10-kpc-wide radial bin, the weighted fitted enclosed mass is

$$M_{\text{fit}}(< r_k, t) = \frac{\sum_{i \in \mathcal{I}_k} N_i M_{\text{sph},i}(< r_k, t)}{\sum_{i \in \mathcal{I}_k} N_i}, \quad (57)$$

where $\mathcal{I}_k$ denotes the selected cohorts whose fitted profiles are evaluated at r_k . We quantify the signed enclosed-mass bias by

$$b_M(t) = \left\langle \frac{M_{\text{fit}}(< r_k, t) - M_{\text{true}}(< r_k, t)}{M_{\text{true}}(< r_k, t)} \right\rangle_k \times 100\%. \quad (58)$$

Figure 4 summarizes the resulting evolution across all four merger ratios. The first panel shows the number of selected cohorts and their mean stellar membership, the second panel shows $\langle C \rangle_w$, and the third panel shows b_M . The distance-matched present-day comparison time is marked by a light-red vertical line.

The number of selected cohorts generally decreases with time in all four merger cases. However, the cohort-count curves do not show a clear monotonic ordering with merger strength. In particular, the stronger mergers do not consistently lose their selected cohorts ear-

lier than the weaker mergers over the full analyzed interval. The compactness generally increases with time, although it does not exhibit a single threshold that separates successful and unsuccessful fits.

The 1:2 merger loses most of its usable cohorts earliest and develops large, rapidly varying mass biases. The 1:5 and 1:10 cases retain selected material for longer, but subsequently exhibit substantial mass fluctuations. The 1:100 merger remains comparatively stable over most of the analyzed interval. Thus, stronger mergers generally shorten the interval over which the fitted mass remains precise.

Occasional late-time snapshots may nevertheless return a small value of b_M . Such isolated agreement does not by itself demonstrate restored reliability. Because the cohorts are selected independently at each snapshot, the late-time sample can be dominated by an atypical subset that happens to remain especially regular. We therefore interpret the time limit of the method from sustained behavior over multiple snapshots, rather than from individual near-zero mass residuals.




Figure 4. Time-dependent performance of the selected-cohort KLD fitting for merger mass ratios 1:2, 1:5, 1:10, and 1:100 in the $30 \leq r/\text{kpc} \leq 40$ radial range. The light-red vertical line marks the distance-matched present-day comparison time at snapshot 055 ($t = 1.406$ Gyr), defined using the fiducial 1:10 merger. **(a)** Number of selected coherent cohorts. Curve color identifies the merger ratio, while the local color saturation encodes the average stellar membership of the selected cohorts; more saturated segments correspond to larger average membership. **(b)** Cohort-member-number-weighted mean six-dimensional phase-space compactness, $\langle C_{6D} \rangle_w$. **(c)** Signed enclosed-mass bias, b_M , of the cohort-member-number-weighted fitted mass profile relative to the true dark-matter enclosed-mass profile measured directly from the snapshot particles and averaged over the radial grid points in the selected range. The background is white when all four merger cases retain at least one selected cohort and becomes progressively darker grey as fewer cases retain usable cohorts. Stronger mergers generally produce larger and more rapidly varying mass biases, whereas the 1:100 merger remains comparatively stable over most of the analyzed interval. The phase-space compactness evolves with time but does not by itself define a universal threshold for successful potential recovery.

4. POST-SELECTION RELIABILITY AND FAILURE DIAGNOSTICS

The selected-cohort results reveal two distinct ways in which a mass estimate can fail after observable coherence cuts have been applied. The KLD optimizer may select a spherical potential different from the cohort-wise targeted potential, and the targeted spherical description may itself differ from the global spherically averaged mass profile. We first separate these two levels alge-

braically. We then examine two non-exclusive reasons for fitted-to-targeted failure: artificial over-clustering of a broadened targeted-action cloud and degeneracy across the NFW parameter plane. Neither mechanism is an additional additive mass-error component; both can operate within the single fitted-to-targeted term.

4.1. Decomposition of the Total Enclosed-mass Bias




Figure 5. Decomposition of the signed enclosed-mass bias for the 1:10 merger in the $20 \leq r/\text{kpc} \leq 30$ radial range. **(a)** Fitted-to-targeted contribution, $b_{\text{fit} \rightarrow \text{tar}}(r, t)$, measuring the difference between the cohort-weighted KLD-best-fit and spherical targeted mass profiles. **(b)** Targeted-to-true contribution, $b_{\text{tar} \rightarrow \text{true}}(r, t)$, measuring the difference between the cohort-weighted spherical targeted profile and the globally spherically averaged BFE mass profile. **(c)** Total fitted-to-true bias, $b_{\text{fit} \rightarrow \text{true}}(r, t)$. At every valid radius–time pixel, panel (c) is the sum of panels (a) and (b). Positive values indicate an overestimate of the enclosed mass, and negative values indicate an underestimate. Gray regions indicate radius–time pixels for which a valid comparison is unavailable. Direct numerical evaluation verifies the additive relation to floating-point precision: the maximum absolute closure residual is 1.510×10^{-14} percentage points and the RMS residual is 1.215×10^{-15} percentage points.

Before examining the action-space mechanisms that can make the KLD fit unreliable, we separate the total enclosed-mass bias into two contributions. The first measures whether the KLD-best-fit spherical potential recovers the spherical targeted potential. The second measures whether that targeted potential reproduces the

globally spherically averaged mass profile of the full BFE field.

Let $M_{\text{fit},w}(< r, t)$ and $M_{\text{tar},w}(< r, t)$ denote, respectively, the cohort-weighted enclosed masses of the KLD-best-fit and spherical targeted potentials. Let $M_{\text{true}}(< r, t)$ denote the globally spherically averaged enclosed mass measured from the full BFE model. We define the

fitted-to-targeted component as

$$b_{\text{fit} \rightarrow \text{tar}}(r, t) = 100 \frac{M_{\text{fit},w}(< r, t) - M_{\text{tar},w}(< r, t)}{M_{\text{true}}(< r, t)}. \quad (59)$$

This term isolates the mass difference introduced when the KLD fit selects a spherical potential different from the targeted potential.

The targeted-to-true component is

$$b_{\text{tar} \rightarrow \text{true}}(r, t) = 100 \frac{M_{\text{tar},w}(< r, t) - M_{\text{true}}(< r, t)}{M_{\text{true}}(< r, t)}. \quad (60)$$

This term remains even if the KLD fit recovers the targeted potential exactly. It measures the difference between the spherical targeted description sampled by the selected cohorts and the globally averaged mass profile of the full BFE field.

The total fitted-to-true mass bias is

$$b_{\text{fit} \rightarrow \text{true}}(r, t) = 100 \frac{M_{\text{fit},w}(< r, t) - M_{\text{true}}(< r, t)}{M_{\text{true}}(< r, t)}. \quad (61)$$

Because all three quantities use the same $M_{\text{true}}(< r, t)$ denominator, they obey the exact additive decomposition

$$b_{\text{fit} \rightarrow \text{true}} = b_{\text{fit} \rightarrow \text{tar}} + b_{\text{tar} \rightarrow \text{true}}. \quad (62)$$

No multiplicative cross-term is required under these definitions.

Figure 5 shows the three terms for the 1:10 merger in the $20 \leq r/\text{kpc} \leq 30$ radial range. The fitted-to-targeted component in panel (a) contains substantial positive and negative structure across radius and time. The largest and most spatially extended deviations become especially prominent after approximately 1.5 Gyr, although localized departures are also present at earlier times. These deviations are the part of the total mass error that can be connected directly to failure of the KLD fit to recover the spherical targeted potential.

The targeted-to-true component in Figure 5(b) is generally smaller and more spatially restricted than the fitted-to-targeted component, but it is not negligible. A particularly coherent negative feature appears around $t \simeq 1.5$ Gyr, while additional localized positive and negative differences occur at later times. These structures

show that the spherical targeted potential does not always reproduce the global BFE enclosed mass even when evaluated over the same selected sample.

The total bias in panel (c) visually resembles the fitted-to-targeted component over much of the available radius–time domain. This indicates that failure to recover the targeted potential is frequently the larger contribution to the total mass-error structure in this example. Nevertheless, panel (b) modifies the total bias wherever the targeted-to-true difference is appreciable. The two signed components can reinforce one another when they have the same sign or partially cancel when they have opposite signs. Consequently, a small total bias at an individual pixel does not necessarily imply that both component errors are small.

The direct closure test summarized in the caption of Figure 5 confirms that the signed decomposition is exact to floating-point precision. This test validates the implementation; it does not introduce a third physical source of mass bias.

The decomposition also establishes the scope of the following diagnostics. The targeted-action spread, KLD over-clustering, KLD-score degeneracy, and median-action degeneracy concern the origin of $b_{\text{fit} \rightarrow \text{tar}}$: they test why the KLD procedure may select a spherical potential different from the targeted potential. They do not directly explain $b_{\text{tar} \rightarrow \text{true}}$, which instead reflects the difference between the spherical targeted description and the globally averaged full BFE mass profile. Moreover, the action-space failure mechanisms are not themselves additive mass-error components and may coexist within the single fitted-to-targeted term.

4.2. Why KLD Maximization Can Miss the Targeted Potential

The KLD method assumes that the cohort-wise targeted potential is represented by the trial model that produces the strongest radial-action clustering. The following diagnostics show why this implication can fail even after the observable selection has removed strongly irregular material.




Figure 6. Relation between the weighted targeted radial-action spread, $S_{J_r,\text{tar},w}$, and the weighted KLD excess, $\Delta D_{\text{KL},w}$, for the 1:10 merger in the $20 \leq r/\text{kpc} \leq 30$ radial range. Each point represents one valid radius-snapshot measurement. The points are colored by snapshot time in panel (a) and by radius in panel (b). Positive $\Delta D_{\text{KL},w}$ indicates that the KLD fit produces a more strongly clustered action cloud than the targeted potential. The distribution broadens toward larger KLD excess as the targeted action cloud becomes more extended, although the relation contains substantial intrinsic scatter. The time coloring in panel (a) shows that earlier snapshots preferentially populate the upper KLD-excess envelope, suggesting that their broad action clouds remain more readily compressible by a mismatched trial potential.




Figure 7. Relation between the weighted targeted radial-action spread, $S_{J_r,\text{tar},w}$, and the signed fitted-to-targeted enclosed-mass bias, $b_{\text{fit} \rightarrow \text{tar}}$, for the 1:10 merger in the $20 \leq r/\text{kpc} \leq 30$ radial range. The points are colored by snapshot time in panel (a) and by radius in panel (b). The range of permitted signed mass biases widens as the targeted action cloud becomes more extended, but the bias can take either sign. Thus, broadening of the targeted action cloud is associated with reduced fitting reliability without uniquely determining the resulting mass error. Earlier snapshots also preferentially occupy the more extreme parts of the mass-bias envelope, although the mass consequence of over-clustering depends on the direction of the fitted displacement in NFW parameter space.

4.2.1. Artificial over-clustering of a broadened targeted-action cloud

The first failure mechanism arises when the selected cohort is not compact in radial action even in the targeted potential. We characterize this broadening using

the weighted targeted-action spread $S_{J_r,\text{tar},w}(r,t)$. We then compare it with the weighted KLD excess

$$\Delta D_{\text{KL},w} = D_{\text{KL},w}(\hat{\theta}) - D_{\text{KL},w}(\theta_{\text{tar}}), \quad (63)$$

where θ_{tar} denotes the targeted-potential parameters and $\hat{\theta}$ denotes the KLD-best-fit parameters. Positive $\Delta D_{\text{KL},w}$ means that the fitted potential produces a more strongly clustered action distribution than the targeted potential.

The current 1:10 result for $20 \leq r/\text{kpc} \leq 30$ is shown in Figure 6. As $S_{J_r,\text{tar},w}$ increases, the occupied region extends toward larger positive $\Delta D_{\text{KL},w}$. The largest KLD excesses occur primarily among the more broadly distributed targeted action clouds, although the substantial scatter shows that targeted-action spread alone does not determine how strongly the fitted potential can over-cluster the distribution.

The time coloring in Figure 6(a) adds a further distinction. Earlier snapshots preferentially populate the upper part of the distribution, while later snapshots with broad targeted-action clouds generally show smaller attainable KLD excesses. A simple physical interpretation is that a broad early-time action cloud can still arise from an orderly stream structure. Changing the trial potential can then move the assigned actions of its members coherently and artificially compress the cloud. After further merger evolution, the stellar kinematics contain more member-dependent disturbance. Changing a smooth spherical potential still remaps every assigned action, but it is less able to undo those unequal perturbations and produce a highly clustered distribution. This interpretation concerns how readily the particular phase-space data can be compressed within the adopted NFW family; it does not imply a universal upper bound on KLD at late times.

The corresponding fitted-to-targeted mass bias is shown in Figure 7. The distribution develops a broader vertical envelope as $S_{J_r,\text{tar},w}$ increases: compact targeted-action clouds are concentrated near small mass bias, whereas broad clouds permit both larger positive and larger negative biases. Earlier snapshots preferentially occupy the more extreme parts of this envelope, consistent with their greater susceptibility to artificial

KLD compression. A large KLD excess gives the optimization a stronger preference for a potential away from the targeted solution, but it does not by itself determine how much the enclosed mass changes.

Targeted-action broadening is therefore a risk indicator rather than a direct predictor of the sign or magnitude of the mass error. Its consequence depends both on how readily the broad cloud can be compressed by the trial-potential family and on the direction in parameter space selected by that compression. The latter dependence is examined in Section 4.2.2.

4.2.2. KLD-score, action-value, and mass-direction degeneracies

Over-clustering is not the only way for KLD maximization to fail. Even when the fitted and targeted action clouds have similar degrees of clustering, the KLD surface need not possess a unique maximum at the targeted parameters. We test this possibility by evaluating the radial actions and their KLD over a two-dimensional grid of trial NFW parameters,

$$\theta = (\rho_0, r_s), \quad (64)$$

centered on the targeted potential. The same cohort and snapshot are used throughout each parameter grid.

Figure 8 shows the resulting KLD landscape for the selected cohort from parent stream 043 in the 1:10 merger at snapshot 55. The cohort has a median radius of $r_{\text{med}} = 25.17 \text{ kpc}$ and was selected from the $20 \leq r/\text{kpc} \leq 30$ radial range. Rather than forming an isolated peak around the targeted potential, the largest KLD values occupy an elongated, approximately one-dimensional region in the (ρ_0, r_s) plane. The near-maximum contour consequently includes multiple combinations of density normalization and scale radius. The KLD score alone cannot strongly distinguish among trial potentials along this direction.




Figure 8. KLD-score degeneracy for a selected cohort from parent stream 043 in the 1:10 merger at snapshot 55. The cohort is selected from $20 \leq r/\text{kpc} \leq 30$ and has median radius $r_{\text{med}} = 25.17 \text{ kpc}$. **(a)** Radial-action KLD evaluated over the trial NFW parameter grid. The axes give the logarithmic offsets of ρ_0 and r_s from their targeted values. The solid contour bounds the region whose KLD is indistinguishable from the grid maximum within the adopted finite-sample score tolerance. The dashed contour bounds trial potentials whose KLD is consistent with that of the targeted potential under the corresponding matched-star comparison. **(b)** Signed enclosed-mass difference between each trial potential and the targeted potential at r_{med} , $b_{\theta \rightarrow \text{tar}}(r_{\text{med}}) = 100[M_{\theta}(< r_{\text{med}}) - M_{\text{tar}}(< r_{\text{med}})]/M_{\text{true}}(< r_{\text{med}})$. The targeted potential and the KLD fit are identified by the markers. The extended, nearly one-dimensional KLD contours demonstrate that similar clustering scores can be produced by different NFW parameter combinations. The KLD tolerances define numerical score equivalence and are not confidence intervals. Gray pixels indicate trial potentials for which the finite-action requirement is not satisfied. The mass effect of a parameter displacement depends on whether it follows or crosses the enclosed-mass contours in panel (b).

The companion mass map in panel (b) shows why an extended KLD ridge does not necessarily imply a large enclosed-mass bias. A change in (ρ_0, r_s) changes the inferred mass only to the extent that the parameter displacement crosses the enclosed-mass contours. A displacement along a nearly constant-mass direction can therefore produce substantially different NFW parameters while retaining a similar enclosed mass at the radius sampled by the cohort.

For the illustrated cohort, portions of the near-maximum and targeted-equivalent KLD regions follow directions that remain close to $b_{\theta \rightarrow \text{tar}}(r_{\text{med}}) = 0$. The KLD score is consequently insufficient to identify unique NFW parameters, but the locally enclosed mass can remain much better constrained than either parameter separately. Conversely, if a KLD-preferred displacement crosses the mass contours, the same parameter non-uniqueness can propagate into a substantial mass bias. The relevant issue is therefore not only the existence of a KLD excess or ridge, but also its orientation relative to the enclosed-mass map.

A second form of non-uniqueness concerns the characteristic action value itself. For every trial potential, we calculate

$$\Delta \text{median } J_r(\boldsymbol{\theta}) = \text{median}_a [J_{r,a}(\boldsymbol{\theta})] - \text{median}_a [J_{r,a}(\boldsymbol{\theta}_{\text{tar}})], \quad (65)$$

where a indexes the stars used at that parameter-grid location. This quantity measures the displacement of the median radial action relative to the targeted potential; it does not measure the width or clustering of the action cloud.

Figure 9 shows that the region producing a targeted-like median action is also extended across the NFW parameter plane. Thus, substantially different combinations of ρ_0 and r_s can assign approximately the same median radial action to the selected cohort. The median radial action is therefore not an injective function of the NFW parameters.




Figure 9. Median-action degeneracy for the same 1:10 cohort, snapshot, and radial range as Figure 8. **(a)** Signed difference between the median radial action in each trial potential and that in the targeted potential, $\Delta \text{median } J_r(\theta)$. The contour marks $|\Delta \text{median } J_r| = 0.10 S_{J_r, \text{tar}}$, where $S_{J_r, \text{tar}}$ is the MAD width of the targeted radial-action cloud. It therefore identifies the boundary of the region whose median action is close to the targeted median under the adopted tolerance. **(b)** Signed enclosed-mass difference $b_{\theta \rightarrow \text{tar}}(r_{\text{med}})$ over the same parameter grid. The extended targeted-median-equivalent region demonstrates that recovery of the targeted median action does not uniquely recover the targeted NFW parameters. Gray pixels indicate trial potentials for which the finite-action requirement is not satisfied.

To control numerical truncation by stars with non-finite radial actions, we define a fixed core sample containing stars with finite actions over at least 75% of the trial-potential grid. A KLD value is retained at a given grid point only when at least 95% of this fixed core remains finite there. A separate map of the fraction of stars with non-finite actions is used to verify which parts of the parameter grid are affected by this truncation. This procedure keeps the stellar sample nearly fixed across the accepted grid while avoiding the removal of the entire map because of a small number of locally non-finite orbits.

The two figures diagnose distinct properties of the action cloud. Figure 8 concerns its degree of clustering, whereas Figure 9 concerns the location of its median radial action. A trial potential can therefore obtain a targeted-equivalent KLD score without reproducing the targeted median, or reproduce the targeted median without being uniquely identified by the KLD objective. Agreement in both scalar quantities would still not prove equality of the complete action distributions on a star-by-star basis.

These results establish the existence of KLD-score and median-action degeneracies for the illustrated cohort, but they also show that parameter identifiability and mass identifiability are different questions. A KLD-preferred displacement can be large in parameter space yet have little effect on the enclosed mass when it fol-

lows a nearly constant-mass direction. Only the component that crosses mass contours changes the inferred mass substantially. The present example demonstrates this directional dependence, but it does not determine how frequently favorable or unfavorable alignments occur across cohorts, snapshots, radial ranges, and merger histories. Over-clustering and degeneracy may also occur simultaneously; they are diagnostic failure routes, not additive components of the mass bias.

The directional result suggests an additional observationally available post-fit screen. For every selected cohort, one may identify the set of trial potentials whose KLD values are statistically or numerically indistinguishable from the maximum and map that entire region to enclosed-mass profiles over the cohort's sampled radial interval. Cohorts for which the near-maximum region spans a narrow mass range may be retained or given greater weight, whereas cohorts permitting a broad mass range should be downweighted or rejected. This criterion preserves cases in which the NFW parameters are non-unique but the local enclosed mass is identifiable. In the NFW examples examined here, adding this preference is compatible with the original cohort selection; a systematic test over a broader potential family remains future work. The screen tests mass identifiability rather than mass accuracy, because an artificially over-clustered near-maximum region can still be displaced as a whole from the targeted solution.

5. SIMULATION-ONLY DIAGNOSTICS OF THE FULL POTENTIAL

The targeted-potential and fitted-to-targeted diagnostics in Section 4 require knowledge of the simulated force field and cannot be reproduced directly from Milky Way observations. We now introduce a second simulation-only comparison that tests the radial-orbit mapping of the fitted spherical model against that of the frozen full BFE field.

5.1. Frozen-snapshot Radial-cycle Diagnostic

At a snapshot t_* , we freeze the BFE potential and integrate a model orbit from each measured stellar phase-space point. If $\tau_{p,k}$ and $\tau_{p,k+1}$ are successive pericentric passages of that frozen orbit, we define

$$\tilde{J}_{r,\text{BFE}}(t_*) \equiv \frac{1}{2\pi} \int_{\tau_{p,k}}^{\tau_{p,k+1}} p_r(\tau) \dot{r}(\tau) d\tau = \frac{1}{2\pi} \int_{\text{radial cycle}} p_r dr. \quad (66)$$

In an integrable spherical potential this construction reduces to the canonical radial action. In the frozen, non-spherical merger potential, however, one radial oscillation need not be a fundamental cycle of an invariant three-torus. The result may depend on phase, cycle

number, and the radial-cycle definition. We therefore treat $\tilde{J}_{r,\text{BFE}}$ as a radial-cycle diagnostic rather than a canonical action.

For cohort i , using only stars with finite values in both calculations, we define

$$\Delta\tilde{J}_{r,\text{med},i} = \text{median}(J_{r,\text{fit},i}) - \text{median}(\tilde{J}_{r,\text{BFE},i}) \quad (67)$$

and

$$\Delta D_{\text{KL},r,i} = D_{\text{KL}}(J_{r,\text{fit},i}) - D_{\text{KL}}(\tilde{J}_{r,\text{BFE},i}). \quad (68)$$

The matched-star-count-weighted diagnostics are

$$\langle \Delta\tilde{J}_{r,\text{med}} \rangle_w = \frac{\sum_i N_{i,\text{match}} \Delta\tilde{J}_{r,\text{med},i}}{\sum_i N_{i,\text{match}}}, \quad (69)$$

$$\langle \Delta D_{\text{KL},r} \rangle_w = \frac{\sum_i N_{i,\text{match}} \Delta D_{\text{KL},r,i}}{\sum_i N_{i,\text{match}}}. \quad (70)$$

These quantities do not supply the theoretical basis of the KLD fit and are not observable inputs. They diagnose disagreement between the radial-orbit structures generated by the fitted spherical model and by the frozen full simulated field.

5.2. Outer-halo Radial-orbit Disagreement




Figure 10. Comparison of the KLD-fitted spherical NFW model with the frozen-snapshot BFE field for merger mass ratios 1:2, 1:5, 1:10, and 1:100 in the $80 \leq r/\text{kpc} \leq 100$ radial range. Only BFE-valid snapshot intervals within $1.41 \leq t/\text{Gyr} \leq 1.76$ contribute. (a) Number of selected coherent cohorts. Color identifies the merger ratio, while the saturation of each curve segment encodes the average stellar membership of the selected cohorts. More saturated segments correspond to larger average membership. (b) Matched-star-count-weighted signed median displacement, $\langle \Delta\tilde{J}_{r,\text{med}} \rangle_w$, scaled by 10^4 as indicated on the axis. (c) Matched-star-count-weighted signed clustering difference, $\langle \Delta D_{\text{KL},r} \rangle_w$. (d) Signed enclosed-mass bias, b_M , of the fitted spherical model relative to the globally spherically averaged BFE mass profile. The action diagnostics are defined in Equations 69 and 70. White background denotes times at which all four merger cases lie within BFE-valid intervals; progressively darker gray denotes coverage by fewer cases. Missing segments occur when the BFE comparison is unavailable or when no selected cohort contributes a finite measurement.

Figure 10 compares the evolution of the selected stream sample, the disagreement between radial actions in the KLD-fitted spherical potential and radial-cycle integrations in the frozen BFE field, and the resulting enclosed-mass bias in the outer-halo range $80 \leq r/\text{kpc} \leq$

100. The action differences and the fitted mass bias exhibit related time dependence. This covariation indicates that disagreement in the radial orbital structure generated by the fitted spherical and frozen full potentials accompanies poor mass recovery. It does not, by

itself, establish that the radial-action difference is an independent or additive source of the mass bias. Rather, it provides a dynamical diagnostic of the mismatch already present between the fitted and true orbital mappings.

The figure is restricted to $1.41 \leq t/\text{Gyr} \leq 1.76$, which provides the largest overlapping BFE-valid coverage across the merger cases. This window does not include the principal BFE-valid interval of every individual simulation. In particular, much of the BFE-valid evolution of the 1:2 merger lies outside the displayed interval, so that case contributes only limited information to the comparison. The apparent differences among merger ratios in this restricted window should therefore be interpreted together with their unequal temporal coverage.

5.3. Dependence on Perturber Orbital Configuration

Figure 11 compares the standard, circular, and radial 1:5 merger configurations over their respective BFE-valid intervals. The radial configuration appears most favorable for the KLD fitting method over the periods analyzed here. Although the close radial encounter may produce a strong transient disturbance, the strongest forcing is plausibly concentrated into a relatively short interval. Outside this interval, coherent cohorts remain or reappear, and the fitted-minus-BFE action discrepancies and enclosed-mass biases are comparatively modest over much of the sampled evolution.

The circular configuration is less favorable in terms of sustained tracer availability. One possible interpretation is that prolonged merger-induced forcing continuously distorts the halo and stream population. The surviving action clouds generally have positive $\Delta D_{\text{KL},r}$ and negative $\Delta \tilde{J}_{r,\text{med}}$, indicating that the optimized spherical NFW model produces an action distribution that is more clustered and shifted toward lower radial actions relative to the BFE field.

This orbital comparison should be regarded as suggestive rather than as a universal ranking of perturber trajectories, since the analysis is restricted to intervals in which the BFE provides a sufficiently reliable representation. However, the comparison indicates that not only merger mass ratio, but also the duration and geometry of the encounter, affect the survival of coherent stream material and the form in which the action-space fitting failure appears.

6. DISCUSSION AND CONCLUSIONS

6.1. What Observable Cohort Selection Can and Cannot Guarantee

The selection procedure developed in Sections 2.5 and 3.2 removes stream material that is visibly folded, phase mixed, or kinematically irregular. A surviving cohort is therefore more likely to have experienced an approximately common targeted-action change, leaving a compact cloud that can support action-clustering inference.

This is the principal reason that the post-selection sample can recover a useful inner-halo mass profile at the distance-matched 1:10 snapshot.

The selection remains a probabilistic dynamical indicator rather than a proof. A cohort may pass the present-day cuts even if its targeted-action cloud was broadened by earlier member-dependent forcing. Conversely, the absence of a cohort at one snapshot does not imply that every material stream has been destroyed; it means that no tested segment then satisfies the adopted observable requirements. Because selection is repeated independently, the cohorts contributing at different snapshots need not be the same.

The remaining failure routes occur at three levels. First, the system may no longer contain enough selected cohorts to constrain the potential. Second, surviving cohorts may fail to recover their spherical targeted potentials because a different trial potential artificially over-clusters the assigned actions or because the KLD surface is degenerate. Third, even exact targeted-potential recovery need not reproduce the globally averaged true mass profile. The exact fitted-to-targeted plus targeted-to-true decomposition prevents these logically distinct effects from being conflated. In the illustrated 1:10 inner-halo case, the targeted-to-true contribution is generally smaller and more spatially restricted, whereas failure of the KLD fit to recover the targeted spherical model supplies most of the extended mass-bias structure.

The time ordering in Figures 6 and 7 admits a simple physical interpretation. At early stages, a stream can be broadened yet retain an orderly relation among its stellar phase-space points. A mismatched smooth potential may then undo part of that orderly variation and create an artificially compact action cloud. At later stages, the kinematic disturbances can become more member dependent, so changing the spherical potential is less able to align all assigned actions simultaneously. This does not impose a smaller mathematical upper bound on late-time KLD; it states that high clustering is less accessible for those more irregular data within the tested NFW family.

The parameter-plane diagnostics add a separate directional condition. KLD overestimation or parameter non-uniqueness changes the inferred mass only through the component of the fitted displacement that crosses enclosed-mass contours. A displacement along a nearly constant-mass direction can leave the locally inferred mass nearly unchanged even when the KLD or NFW parameters differ from their targeted values. The post-fit screen proposed in Section 4.2.2 therefore evaluates the enclosed-mass range of the full near-maximum KLD region rather than requiring a unique parameter maximum.

Several validation quantities used here require simulation knowledge: the true time-dependent force, the cohort-wise targeted-potential sequence, targeted ac-




Figure 11. Comparison of the KLD-fitted spherical NFW model with the frozen-snapshot BFE field for the standard, circular, and radial 1:5 merger configurations in the $80 \leq r/\text{kpc} \leq 100$ radial range. Only BFE-valid snapshot intervals within $1.41 \leq t/\text{Gyr} \leq 1.76$ contribute to the comparison. **(a)** Number of selected coherent cohorts. Curve color identifies the orbital configuration, while increasing color saturation represents increasing average stellar membership per cohort. **(b)** Matched-star-count-weighted signed median displacement, $\langle \Delta \tilde{J}_{r,\text{med}} \rangle_w$, scaled by 10^4 as indicated on the axis. **(c)** Matched-star-count-weighted signed clustering difference, $\langle \Delta D_{\text{KL},r} \rangle_w$. **(d)** Signed enclosed-mass bias, b_M , of the fitted spherical model relative to the globally spherically averaged BFE mass profile. White background indicates times at which all three configurations lie within BFE-valid intervals; progressively darker gray indicates that fewer configurations provide a valid BFE comparison. Missing curve segments indicate that no valid diagnostic is available.

tions, the direct dark-matter enclosed mass, and the frozen-BFE radial-cycle integrations. In Galactic data, reliability must instead be assessed through observable coherence, resampling stability, agreement among independently fitted cohorts, consistency across radial ranges or action components, sensitivity to fitting bounds, and the mass range spanned by the near-maximum KLD region.

6.2. Future Potential Families

The spherical NFW family is deliberately economical and keeps the actions well defined and computationally tractable, but it cannot represent the full angular structure of a merger-distorted halo. A natural next step is to repeat the analysis with an axisymmetric or triaxial action-computable family. Such a model could reduce the targeted-to-true mismatch and test whether some fitted-to-targeted failures currently attributed to over-clustering instead reflect insufficient flexibility of the spherical family. The extension must preserve a well-defined action construction—for example through an integrable or controlled approximately integrable model—because a generic triaxial potential is not automatically integrable. It will also require more stellar coverage and stronger regularization to prevent the added parameters from worsening the KLD degeneracies diagnosed here.

6.3. Conclusions

Our principal conclusions are as follows:

1. A time-dependent non-integrable merger field can be analyzed through cohort-specific sequences of instantaneously frozen spherical targeted poten-

tials without treating the true potential as integrable.

2. The assigned targeted action evolves through an exactly derived residual-force term and a targeted-map evolution term. Observable cohort selection is designed to favor material for which the member-to-member differences in their accumulated sum remain small.
3. Selected cohorts recover a useful inner-halo enclosed-mass profile at the distance-matched 1:10 snapshot, but stronger mergers generally shorten the interval of stable recovery and eventually reduce the supply of usable material.
4. Broadened targeted-action clouds are more vulnerable to artificial KLD compression, although later, more irregular kinematics can be less readily remapped into a highly clustered cloud. Targeted-action spread is therefore a risk indicator, not a deterministic predictor of mass-bias sign or magnitude.
5. KLD-score and median-action degeneracies need not imply a large mass error. Their consequence depends on the orientation of the KLD-preferred region relative to enclosed-mass contours, motivating a post-fit mass-identifiability screen.
6. Observable selection improves robustness but cannot guarantee unique potential recovery or an unbiased global mass profile. Simulation-only BFE comparisons remain valuable for calibrating these

limitations and for guiding future applications to Galactic data.

REFERENCES

- Arora, A., Sanderson, R. E., Panithanpaisal, N., et al. 2022, *ApJ*, 939, 2, doi: [10.3847/1538-4357/ac93fb](https://doi.org/10.3847/1538-4357/ac93fb)
- Besla, G., Kallivayalil, N., Hernquist, L., et al. 2007, *ApJ*, 668, 949, doi: [10.1086/521385](https://doi.org/10.1086/521385)
- Brooks, R. A. N., Garavito-Camargo, N., Johnston, K. V., et al. 2025, *ApJ*, 978, 79, doi: [10.3847/1538-4357/ad93a7](https://doi.org/10.3847/1538-4357/ad93a7)
- Brooks, R. A. N., Sanders, J. L., Lilleengen, S., Petersen, M. S., & Pontzen, A. 2024, *MNRAS*, 532, 2657, doi: [10.1093/mnras/stae1565](https://doi.org/10.1093/mnras/stae1565)
- Erkal, D., Li, T. S., Koposov, S. E., et al. 2018, *MNRAS*, 481, 3148, doi: [10.1093/mnras/sty2518](https://doi.org/10.1093/mnras/sty2518)
- Ferdosi, B. J., Buddelmeijer, H., Trager, S. C., Wilkinson, M. H. F., & Roerdink, J. B. T. M. 2011, *A&A*, 531, A114, doi: [10.1051/0004-6361/201116878](https://doi.org/10.1051/0004-6361/201116878)
- Garavito-Camargo, N., Besla, G., Laporte, C. F. P., et al. 2019, *ApJ*, 884, 51, doi: [10.3847/1538-4357/ab32eb](https://doi.org/10.3847/1538-4357/ab32eb)
- Kallivayalil, N., van der Marel, R. P., Besla, G., Anderson, J., & Alcock, C. 2013, *ApJ*, 764, 161, doi: [10.1088/0004-637X/764/2/161](https://doi.org/10.1088/0004-637X/764/2/161)
- Petersen, M. S., & Peñarrubia, J. 2020, *MNRAS*, 494, L11, doi: [10.1093/mnras/slz029](https://doi.org/10.1093/mnras/slz029)
- Reino, S., Rossi, E. M., Sanderson, R. E., et al. 2021, *MNRAS*, 502, 4170, doi: [10.1093/mnras/stab304](https://doi.org/10.1093/mnras/stab304)
- Reino, S., Sanderson, R. E., Panithanpaisal, N., Rossi, E. M., & Kuijken, K. 2022, *MNRAS*, 509, 5365, doi: [10.1093/mnras/stab3176](https://doi.org/10.1093/mnras/stab3176)
- Sanders, J. L., & Binney, J. 2016, *MNRAS*, 457, 2107, doi: [10.1093/mnras/stw106](https://doi.org/10.1093/mnras/stw106)
- Sanderson, R. E., Helmi, A., & Hogg, D. W. 2015, *ApJ*, 801, 98, doi: [10.1088/0004-637X/801/2/98](https://doi.org/10.1088/0004-637X/801/2/98)
- Shipp, N., Li, T. S., Pace, A. B., et al. 2019, *ApJ*, 885, 3, doi: [10.3847/1538-4357/ab44bf](https://doi.org/10.3847/1538-4357/ab44bf)
- Vera-Ciro, C., & Helmi, A. 2013, *ApJL*, 773, L4, doi: [10.1088/2041-8205/773/1/L4](https://doi.org/10.1088/2041-8205/773/1/L4)
- Wu, X., Famaey, B., Gentile, G., Perets, H., & Zhao, H. 2008, *MNRAS*, 386, 2199, doi: [10.1111/j.1365-2966.2008.13198.x](https://doi.org/10.1111/j.1365-2966.2008.13198.x)